

Eric_Hirsch_622_Final_Assignment

Predicting the Space Titanic Kaggle Competition

Eric Hirsch

12/11/2022

Contents

Discussion	2
Introduction	2
Prediction using the Kaggle Spaceship Titanic Data Set	3
The Business Problem	3
Data Summary and Preparation	3
Missing Values	4
Outliers	4
Feature Engineering	4
Modelling	4
Choosing and Testing Models	4
Hyperparameter Tuning	5
Results	5
Discussion	5
Code	7
1. Data Exploration	7
A. Summary Statistics	7
B. Distributions: Skewedness and Outliers	8
C. Multicollinearity	16
D. Missing Values	17
E. First Pass Logistic Regression: 9th percentile	22
Data Preparation and Feature Engineering	24
1. Create groups based on the Passenger ID	24
2. Create Cabin Variables	26
3. Create Dummy Variables	26
4. Implement Interaction Features	27

5. Perform Logistic Regression With Engineered Features: 26th Percentile	27
More Complex Models	30
1. Perform SVM: 70th Percentile	30
2. Perform limited Neural Networks	34
3. Tree Algorithms 1: Perform Random Forest: 34th Percentile	36
4. Tree Algorithms 2: Perform XGBoost Untuned: 73rd Percentile	38
Hypertuning	39
1. Tune XGBoost: 78th Percentile	39
2. Tune SVM Radial: No Improvement	42
Discussion	43

```
knitr::opts_chunk$set(echo = TRUE, warning = FALSE, message = FALSE)
```

```
dfTable <- read.csv("D:\\Rstudio\\Cuny_622\\Space\\622a.csv")
```

Discussion

Introduction

In machine learning, we predict target variables based on input variables. For this final exercise, we will apply various machine learning algorithms to a Kaggle data set (Spaceship Titanic) in order to predict which passengers have been transported to another dimension.

While it's tempting to throw as many algorithms at the problem as possible to see what sticks, the statistical fact is that while it is rare that a poor model will perform well on a holdout set, the chances of making false conclusions based on performance increases if we simply try one model after another. Besides, if we don't understand our model and our data, and the model becomes much more difficult to troubleshoot and maintain.

When choosing models, we are balancing simplicity and complexity, and therefore tendencies to underfit or overfit. When the relationships in the data are simple and certain statistical conditions are met, parametric methods like OLS work well and have the advantage of being easily interpretable. If, for example, we are predicting height from weight, the relationship is simple enough that we can create a linear regression model and capture most of the variation that can be explained for these two variables.

When we increase our dimensions and/or complexity of relationships within the dataset, parametric methods are likely to underfit the data. Even in our simple height and weight example, if the relationship between height and weight varies considerably at lower weights, medium weights and higher weights, spline regression or another nonparametric technique will be necessary. As dimensions and complexity increases, we adopt techniques that are more powerful at morphing the data shape so that we can model the underlying structure, such as trees, SVM and neural nets.

Choosing the more complex algorithm will likely fit the training data better, but may be less interpretable and more subject to overfitting. With this in mind, each of these techniques has its advantages and disadvantages. In my experience with earlier datasets in this class, trees will pick up autonomous clusters in the data set better than SVMs. For example, if there were a small but significant anomalous cluster of individuals for whom height and weight were inversely related, trees will incorporate the cluster while SVMs will ignore it. Of course, clusters like this might signal a missing variable, but not all of the necessary variables will be found in any given data set. On the other hand, when the relationships are more systematic and class

boundaries are clear, SVMs may perform better because the kernel trick allows SVMs to radically change the data shape in order to find the class boundary. SVMs can also perform better when there is less data.

While decision trees are powerful, they are generally more so when bagged (e.g., Random Forest) or boosted (e.g. xgBoost). Because xgBoost is an active learner, it will often have the upper hand in fitting the training data.

One of the biggest advantages of neural networks is that they effectively do the feature engineering for you if you can apply enough layers. They are also subject to the “double descent” phenomenon, which helps with managing underfitting. However, for a student using a home computer like myself, it’s often impractical to take advantage of these facts as the algorithm would run too long. Neural networks, like SVMs, also powerfully change the data shape in order to find class boundaries.

Accurate prediction depends not only on algorithm choice. We also need to engineer features (except possibly in very large neural nets) and tune hyperparameters. We also need to choose metrics that tell us whether or not our model is effective.

Prediction using the Kaggle Spaceship Titanic Data Set

For this exercise I’ve chosen a Kaggle Competition – the Kaggle Spaceship data set. The advantages of using this a competition data set are that we can compare our performance to those of others. Achieving 90% on a holdout set in and of itself tells us nothing - we don’t know if achieving 95% would have been easy or impossible. In this competition, the 2,000 or so submitted accuracies on the leaderboard range from about 76% to 82%, which can give us a good idea of how well our model is working.

The main disadvantages of this data set are that the data is made up and the scenario a bit far-fetched. However, I wanted a data set that had a simple class as a target, as opposed to an image example. The standard Titanic data set has been over analyzed, so this was one of the few good choices left.

The Business Problem

In the year 2912, the Spaceship Titanic, an interstellar passenger liner with almost 13,000 passengers on board, collided with a spacetime anomaly hidden within a dust cloud. Though the ship stayed intact, almost half of the passengers were transported to an alternate dimension. Our job is to predict which passengers were transported by the anomaly using a set of partial records recovered from the spaceship’s damaged computer system.

Data Summary and Preparation

The data set consists of 8693 records and 13 variables, including spending on the ship’s various amenities (VR Deck, Spa, Room Service, Food Court, and Shopping Mall), cabin number, whether the individual was traveling with the group, whether the individual was a VIP, planet of origin and destination, and so on. These columns map to some degree with the original Titanic database.

The target variable, Transported, is roughly equally distributed between false (4315) and true (4378).

Some of the numeric variables, particularly spending variables, are highly skewed - most passengers spend no money while a few spend a great deal. We can see that spending on luxuries (the spa, room service, etc.) is strongly negatively correlated with being transported - this supports the supposition that the rich were spared. Spending on more budget-friendly amenities like the food court and shopping mall are also negatively correlated but less so. Age has a small negative correlation as well.

We decide not to log transform the numeric variables as normal distributions for predictors are not required by our models and interpretability suffers.

Missing Values 1073, or 12%, of records have missing values. The vast majority of missing values are found in the amenity expenditure columns. Oddly, the amenity expenditure rows with missing values are completely independent of each other - there are almost no records where more than one of these values is missing.

This may be an artifact of the fact that the data is manufactured. In order to confirm that there is no systematic relationship between missing data and the target variable, we look at the Chi square between the target and a flag designating missing data. We do this for each amenity expenditure column and find no relationship between missing data and the target variable. We therefore eliminate rows with missing values for the training set. For the test set, we impute the median.

There is very little, even surprisingly little, multicollinearity in the database. In the case of variables that track spending on amenities this is most surprising, and may suggest that passengers were working within a budget and only spent money on the activities they liked most.

Outliers All of the spending variables are highly skewed, with very large instances occurring at the very end of the distribution. However, as most of our techniques are robust for outliers, records with extreme values remain in the database, as there is no reason to think that the spending is a data entry error or an anomalous occurrence.

Feature Engineering The data set holds a number of opportunities for feature engineering. Through testing, it was found that the following new features were significant in predicting transportation. They are:

1. Groups - Passenger ID numbers suggest that some passengers boarded the ship as part of a group. Although we could not establish that members of a group tended to meet the same fate, being a member of a group influenced whether a passenger was transported.
2. Cabins - Cabin codes were parsed for location on the ship - this information was correlated with the target variable.
3. Interaction variables - there were a number of interactions among variables which appeared when the variables were examined in isolation. We only retained one (group passengers were more likely to be transported if they shopped at the mall) because the others did not significantly increase accuracy when running the overall model - but there would be more to explore here.

At this point, a picture of the transported passengers begins to emerge - they are the poorer passengers, most likely to shop on a budget or, even cheaper, spend the trip in cryosleep. They tend to enter the ship in groups and inhabit lower class cabins.

Modelling

Choosing and Testing Models Understanding what models are doing and how is a key part of prediction. We compared tree models, svm, neural net, and logistic regression.

The first task is to understand the requirements of the business problem. In this case, we have a partial roster of passengers where we know who was transported and who wasn't. As for the rest of the passengers, it is our job to predict which of them were transported as well. Insight into the data is necessary insofar as it helps us make better predictions, but we need no more from the data than that. Accuracy is our metric, as this is the metric used in the Kaggle (e.g., there is no penalty for false positives or false negatives which might suggest a different metric). This suggests we should use the most complex model we can that has the highest accuracy.

We used tenfold cross validation and compared models on the metric of accuracy - our results varied widely. There was some, but not perfect, match between the accuracy predicted by the cross validation, and the

accuracy achieved in the Kaggle. For example, going on cross validation alone, the svm with the radial kernel underperformed compared to a linear and polynomial kernel - however, in the Kaggle it significantly outperformed both kernels. Of course, without the Kaggle we would not have known this and would not have chosen the radial kernel to put into production.

On the other hand, our best model on cross validation (xgboost) also performed best on the Kaggle. We suspect the neural net would have performed well also, but we did not have the computer power to hypertune it properly.

Hyperparameter Tuning The caret package in R automatically hypertunes models on basic parameters and chooses the most optimal based on the metric determined by the user. With the exception of xgBoost, all of our models were tuned that way. This left two tasks – tune xgBoost, and experiment with manual tuning of other models.

The tuned xgBoost model (number of rounds was reduce from 55 to 14) increased accuracy from 79% to 80%, and improved our percentile position in the Kaggle from the 73rd to the 78th percentile.

The fact that our radial-kernal SVM had lower cross validation accuracy but higher Kaggle accuracy might have a number of different causes - but one is the possibility that there are idiosyncrasies in the training set that are not found in the Kaggle set. In this case, we can increase sigma to decrease the risk of overfitting.

As we experimented by increasing sigma and decreasing sigma on the radial SVM model. Reducing sigma (tighter fit) improved accuracy during cross validation, but not when applied to the Kaggle. Because the differences were very small (.0026 improvement in accuracy), the models probably only varied by a handful of predictions. Increasing sigma (looser fit) reduced accuracy both in cross validation and on the Kaggle set.

Results

```
knitr::kable(dfTable)
```

Model	Holdout.Set.Accuracy	Kaggle.Accuracy	Kaggle.Percentile
XGBoost Tuned (14 Rounds)	0.8000	0.80383	78.2%
XGBoost Untuned (55 Rounds)	0.7900	0.80243	73.8%
SVM Rad	0.7715	0.80149	70.0%
SVM Rad Lower Sigma	0.7741	0.80032	61.1%
SVM Rad Higher Sigma	0.7669	0.79682	57.3%
SVM Linear	0.7800	0.79775	56.8%
SVM Poly	0.7800	0.79611	55.3%
Random Forest	0.7900	0.78793	34.5%
Logistic Regression with feature engineering	0.7800	0.77975	26.2%
Neural Net (severely limited due to computer power)	0.7600	0.77881	25.5%
Logistic Regression without feature engineering	0.7800	0.69227	9.5%

Discussion

For this exercise, we entered a Kaggle competition (Space Titanic) in order to make predictions and test their accuracy against other competition participants. The target variable was whether or not a passenger was transported to another dimension while the ship was moving through a dangerous and destructive space anomaly. We reached the 78th percentile, with an idea of further modifications which might help us do better.

In a sense, measuring performance in a Kaggle is a cheat - after a time the “unknown” data in the Kaggle test set becomes increasingly “known” as you test more and more models against it. However, it also provides a real-world check on how well your model is actually performing.

After an initial logistic regression, we improved performance in three ways:

1. Feature Engineering
2. Model Selection
3. Hypertuning

The exercise highlighted the importance of Understanding the data in solid and dependable prediction. First, we needed to clean the data of anomalies (missing values, outliers, etc.). In our case, missing values appear to be at random and outliers didn’t present a problem.

Second, an understanding of the data highly conditioned our ability to perform feature engineering. Transported passengers tended to spend their money at the food court rather than get room service, they occupied lower-class cabins, and many of them were in cryosleep during the journey. Knowing this helped us parse out cabin locations and discover interaction terms within the data.

All of our models came in between 77% and 80% accuracy (based on a no information rate of 51%) after tenfold cross validation and optimal hypertuning. This was also the range for individuals who submitted predictions in the Kaggle competition - and so a score of 77% achieved a much lower percentile in the competition than a score a few percentage points higher. Without the Kaggle, we still would have chosen the xgBoost model, but if we were going to choose an SVM we would likely have chosen a less than optimal kernel.

While it makes sense that the boosted tree model outperforms the bagged tree model, it more difficult to speculate on why tree models perform better than the Support Vector Machines. Given the small difference in accuracy between them (.003), it is possible that there is no significant difference, and that with slightly different feature engineering the SVM would be the better model.

There would be a number of ways to improve accuracy even further:

1. Search for more features, particularly interaction features between cryosleep and spending, and understanding better how groupings and cabins work.
2. Borrow a more powerful computer to perform a properly tuned neural net, and then create a prediction set based on neural net, xgBoost and radial SVM.
3. Experiment more with manual tuning of models.

In the meantime, however, we are satisfied with the nearly 80th percentile compared to the 26th percentile using logistical regression which we would have achieved prior to this class.

```
#dfTrain <- read.csv("C:\\Users\\erico\\Documents\\R\\Space\\train.csv")
#dfTest <- read.csv("C:\\Users\\erico\\Documents\\R\\Space\\test.csv")

dfTrain <- read.csv("D:\\Rstudio\\Cuny_622\\Space\\train.csv")
dfTest <- read.csv("D:\\Rstudio\\Cuny_622\\Space\\test.csv")

dfTrain$Transported = ifelse(dfTrain$Transported=="True",1,0)

#dfTrain <- read.csv("C:\\Users\\erico\\Documents\\R\\Space\\train.csv")
#dfTest <- read.csv("C:\\Users\\erico\\Documents\\R\\Space\\test.csv")
```

Code

1. Data Exploration

A. Summary Statistics

The data consists of 8693 records and 14 variables (6 numeric and 8 character). There are a number of missing values and what appear to be skewed distributions among the numeric variables.

```
summary(dfTrain)
```

```
## PassengerId      HomePlanet      CryoSleep      Cabin
## Length:8693      Length:8693      Length:8693      Length:8693
## Class :character  Class :character  Class :character  Class :character
## Mode  :character  Mode  :character  Mode  :character  Mode  :character
##
##
##
## Destination      Age      VIP      RoomService
## Length:8693      Min.   : 0.00  Length:8693      Min.   : 0.0
## Class :character  1st Qu.:19.00  Class :character  1st Qu.: 0.0
## Mode  :character  Median :27.00  Mode  :character  Median : 0.0
##                      Mean  :28.83                Mean  : 224.7
##                      3rd Qu.:38.00                3rd Qu.: 47.0
##                      Max.   :79.00                Max.   :14327.0
##                      NA's   :179                  NA's   :181
## FoodCourt      ShoppingMall      Spa      VRDeck
## Min.   : 0.0  Min.   : 0.0  Min.   : 0.0  Min.   : 0.0
## 1st Qu.: 0.0  1st Qu.: 0.0  1st Qu.: 0.0  1st Qu.: 0.0
## Median : 0.0  Median : 0.0  Median : 0.0  Median : 0.0
## Mean   : 458.1 Mean   : 173.7 Mean   : 311.1 Mean   : 304.9
## 3rd Qu.: 76.0  3rd Qu.: 27.0  3rd Qu.: 59.0  3rd Qu.: 46.0
## Max.   :29813.0 Max.   :23492.0 Max.   :22408.0 Max.   :24133.0
## NA's   :183    NA's   :208    NA's   :183    NA's   :188
## Name      Transported
## Length:8693 Min.   :0.0000
## Class :character 1st Qu.:0.0000
## Mode  :character Median :1.0000
##                      Mean  :0.5036
##                      3rd Qu.:1.0000
##                      Max.   :1.0000
##
```

```
str(dfTrain)
```

```
## 'data.frame': 8693 obs. of 14 variables:
## $ PassengerId : chr "0001_01" "0002_01" "0003_01" "0003_02" ...
## $ HomePlanet : chr "Europa" "Earth" "Europa" "Europa" ...
## $ CryoSleep : chr "False" "False" "False" "False" ...
## $ Cabin : chr "B/O/P" "F/O/S" "A/O/S" "A/O/S" ...
## $ Destination : chr "TRAPPIST-1e" "TRAPPIST-1e" "TRAPPIST-1e" "TRAPPIST-1e" ...
```

```
## $ Age      : num  39 24 58 33 16 44 26 28 35 14 ...
## $ VIP      : chr   "False" "False" "True" "False" ...
## $ RoomService : num  0 109 43 0 303 0 42 0 0 0 ...
## $ FoodCourt  : num  0 9 3576 1283 70 ...
## $ ShoppingMall: num  0 25 0 371 151 0 3 0 17 0 ...
## $ Spa       : num  0 549 6715 3329 565 ...
## $ VRDeck    : num  0 44 49 193 2 0 0 NA 0 0 ...
## $ Name      : chr   "Maham Ofracculy" "Juanna Vines" "Altark Susent" "Solam Susent" ...
## $ Transported : num  0 1 0 0 1 1 1 1 1 1 ...
```

```
dfTrain %>%
  count(dfTrain$Transported)
```

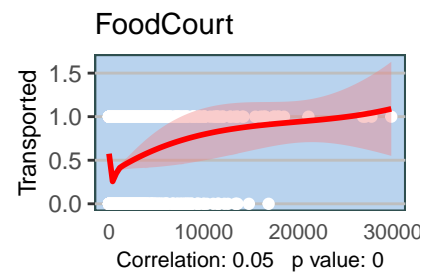
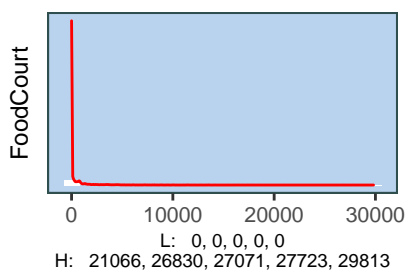
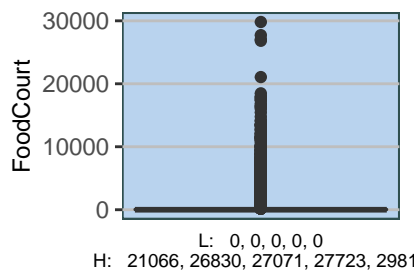
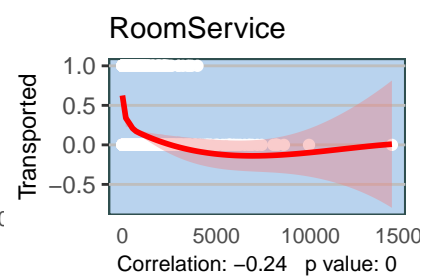
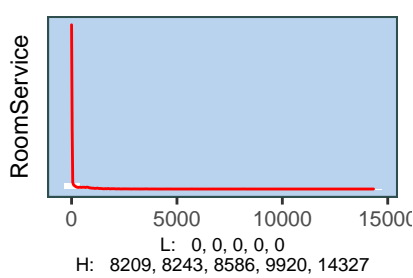
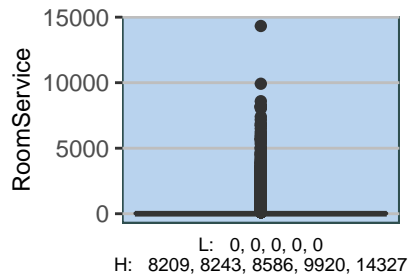
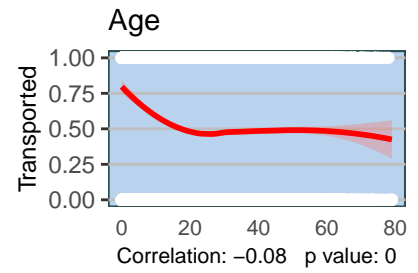
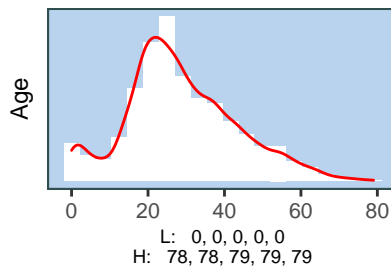
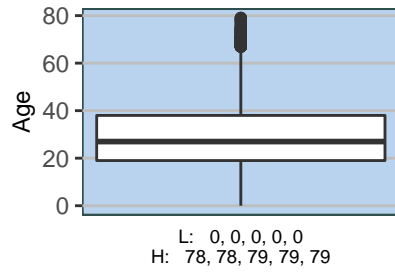
```
##   dfTrain$Transported    n
## 1                    0 4315
## 2                    1 4378
```

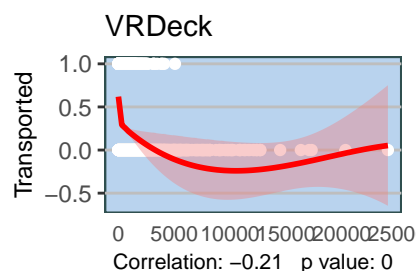
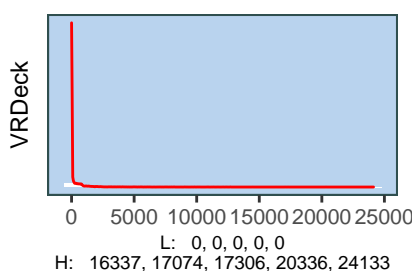
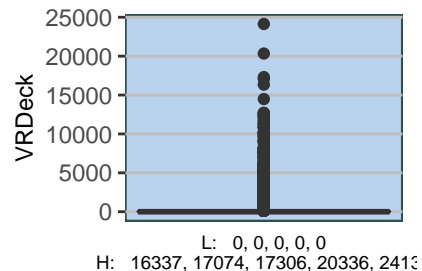
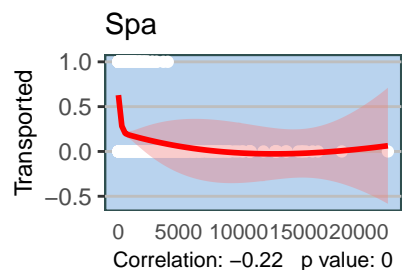
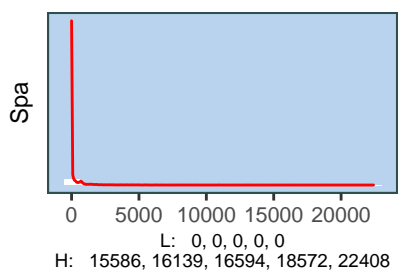
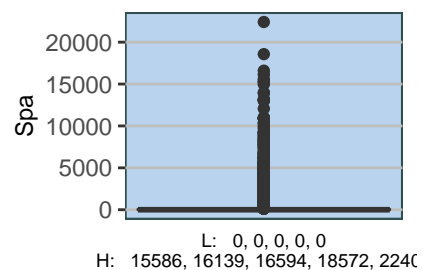
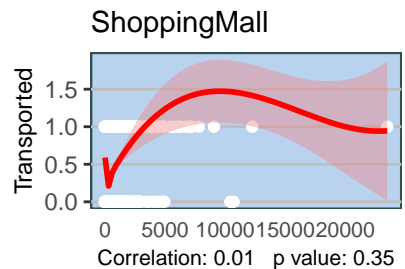
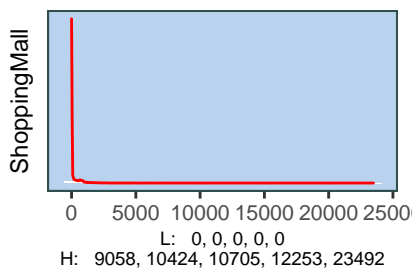
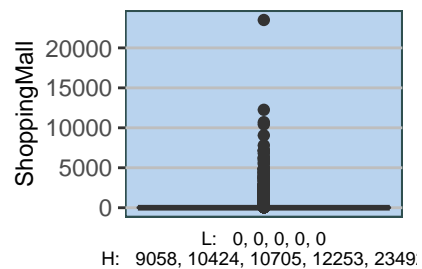
B. Distributions: Skewedness and Outliers

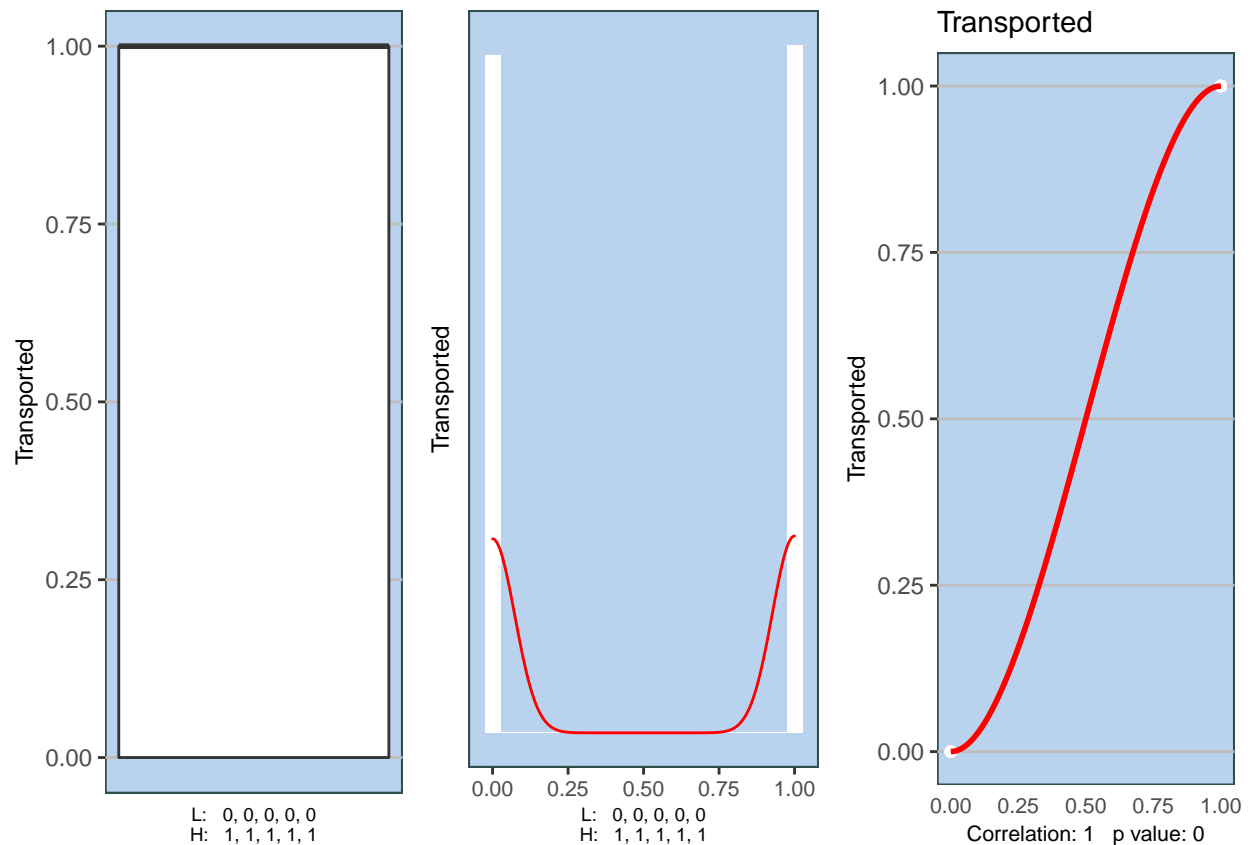
We look at the distribution of all numeric variables. Spending variables are highly skewed - most passengers spend no money while a few spend a great deal. We can see that spending on luxuries (the spa, room service, etc.) is strongly negatively correlated with being transported - this supports the supposition that the rich were spared. Spending on more popular amenities like the food court and shopping mall are also negatively correlated but less so. Age has a small negative correlation as well.

We decide not to log transform the numeric variables as normal distributions for predictors are not required by our models and interpretability suffers.

```
dfTrainD <- dfTrain
EHSummarize_StandardPlots(dfTrainD, "Transported")
```



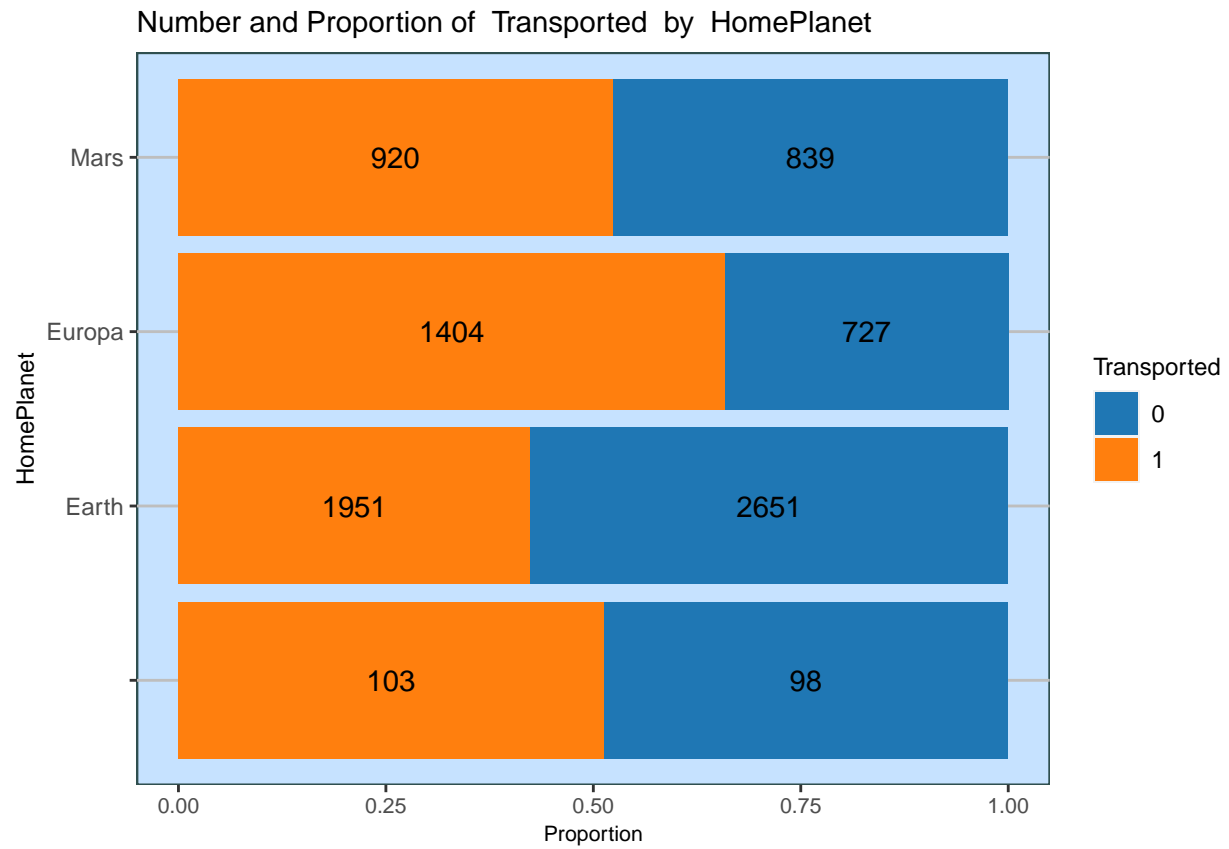


Here we look at count plots for character variables. Home and destination have an association with transported, but cryoSleep is especially important - over 75% of those in cryosleep (presumably a low budget option as food, drink and space are limited) were transported.

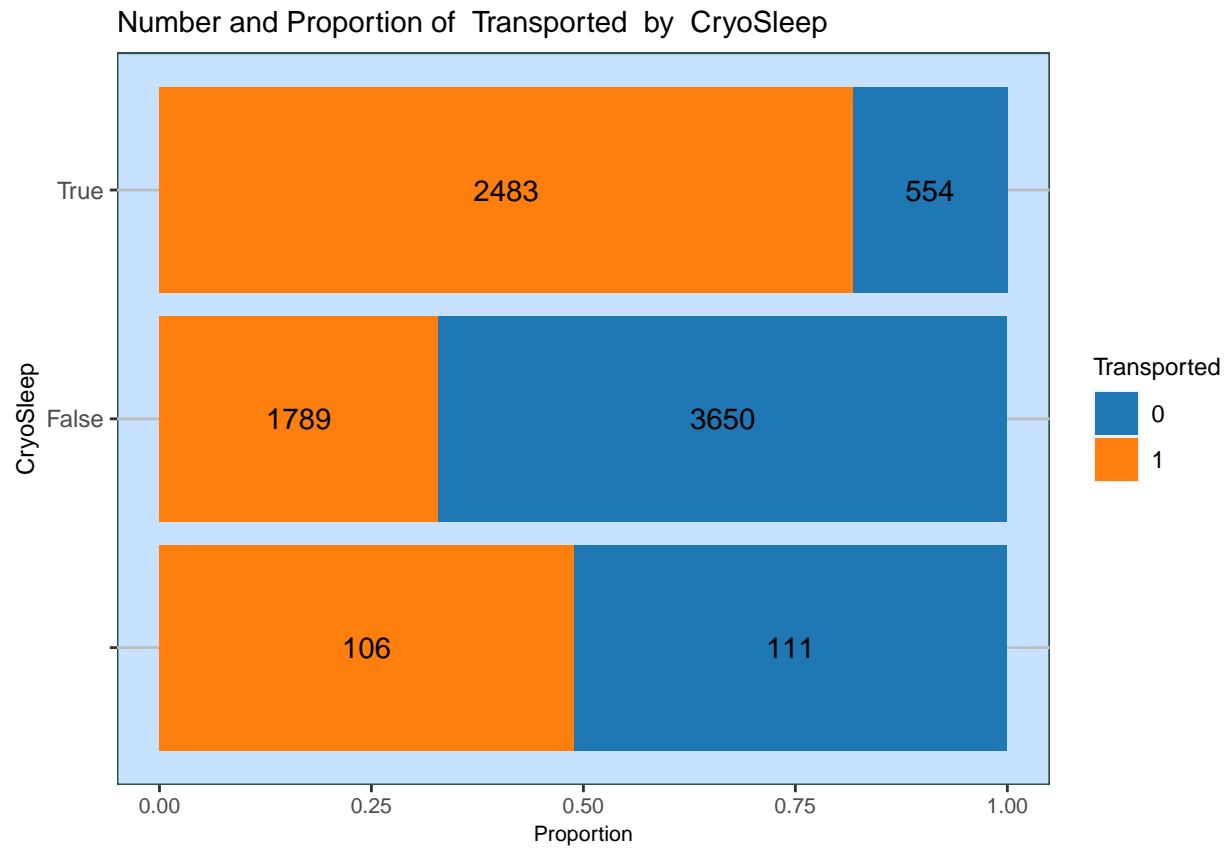
```
dfCount <- dfTrainD %>%
  dplyr::select(HomePlanet, CryoSleep, VIP, Destination, Transported)

dfCount$Transported <- as.character(dfCount$Transported)
EHExplore_TwoCategoricalColumns_Barcharts(dfCount, "Transported")
```

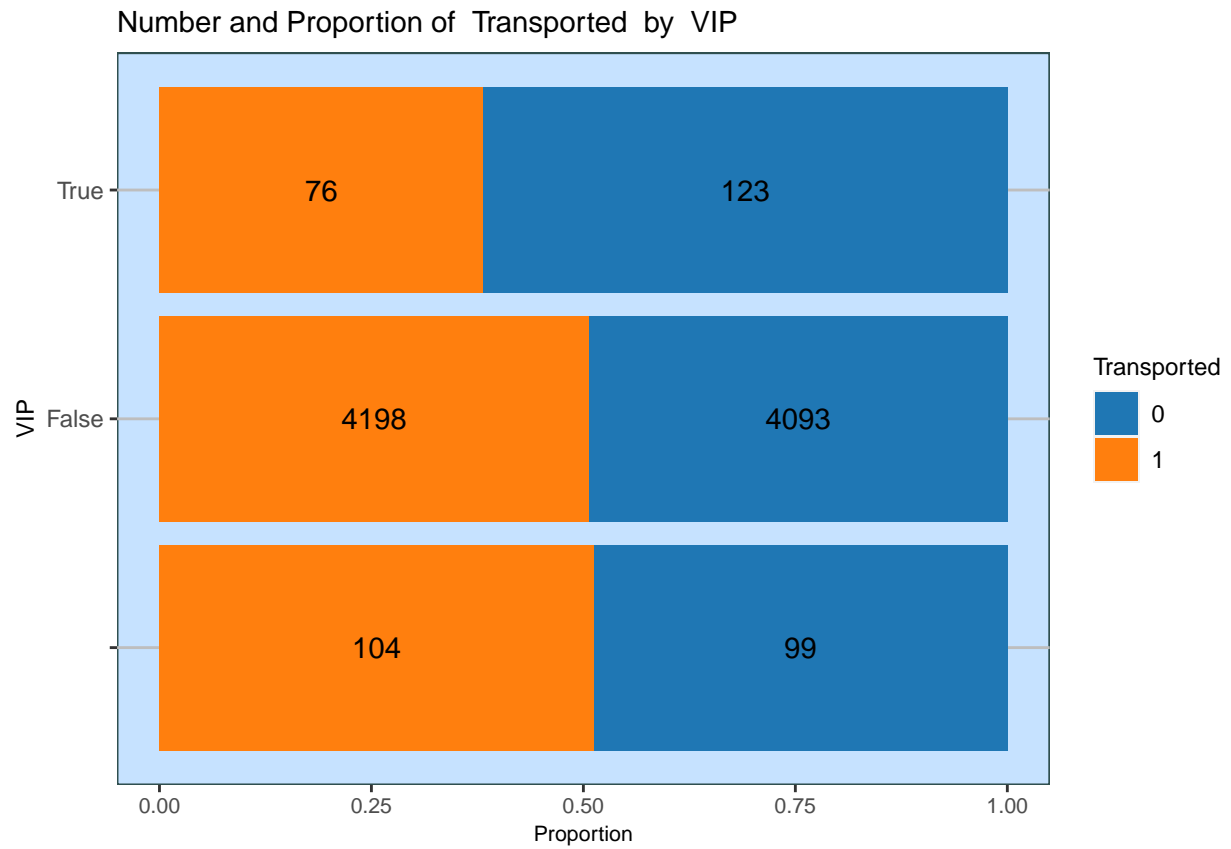
```
## [[1]]
```



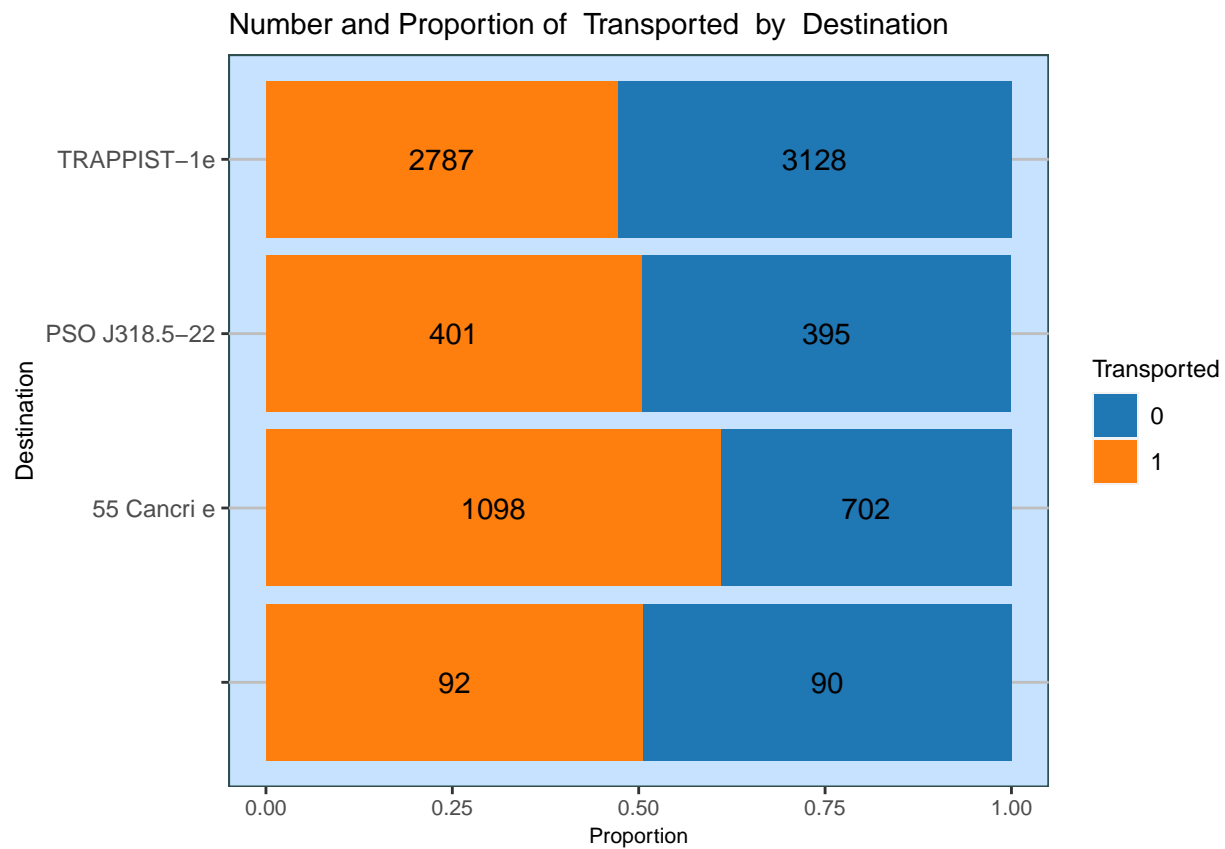
```
##  
## [[2]]
```



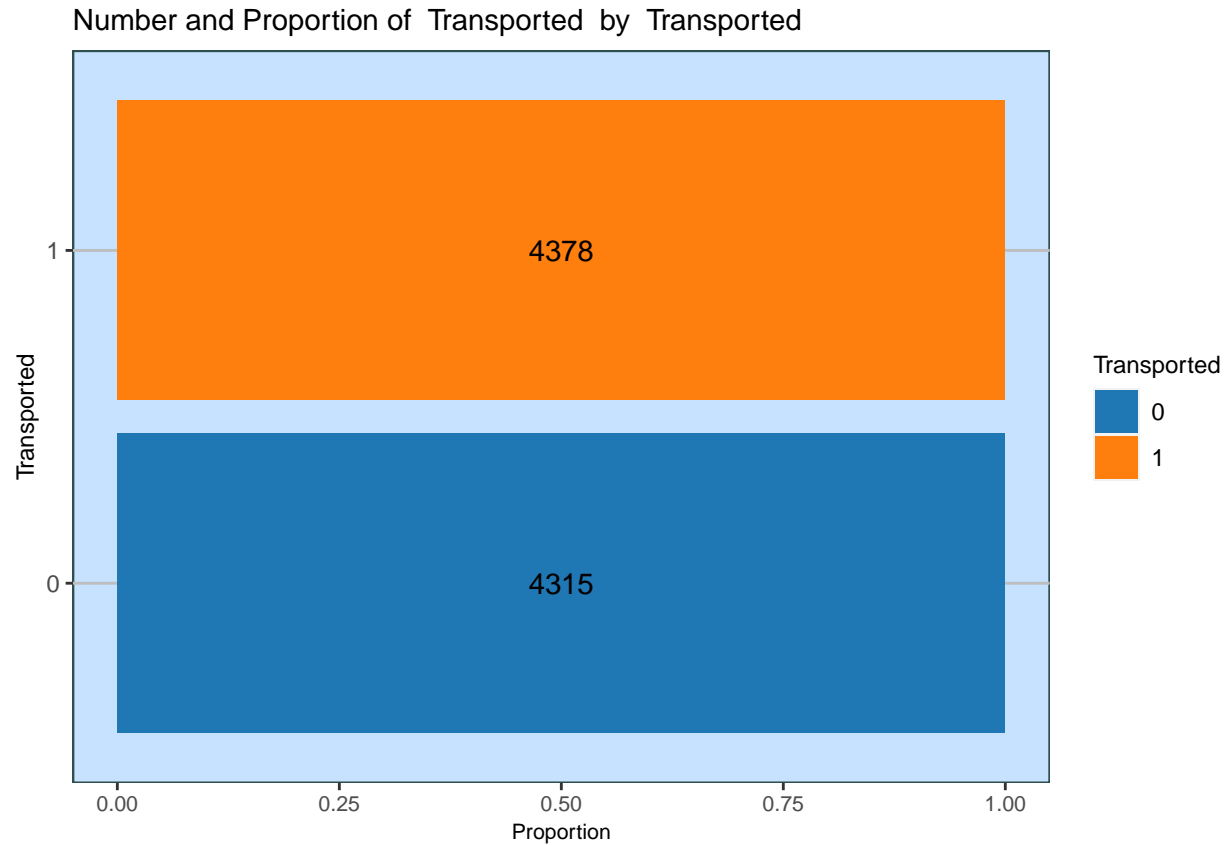
```
##  
## [[3]]
```



```
##  
## [[4]]
```



```
##  
## [[5]]
```



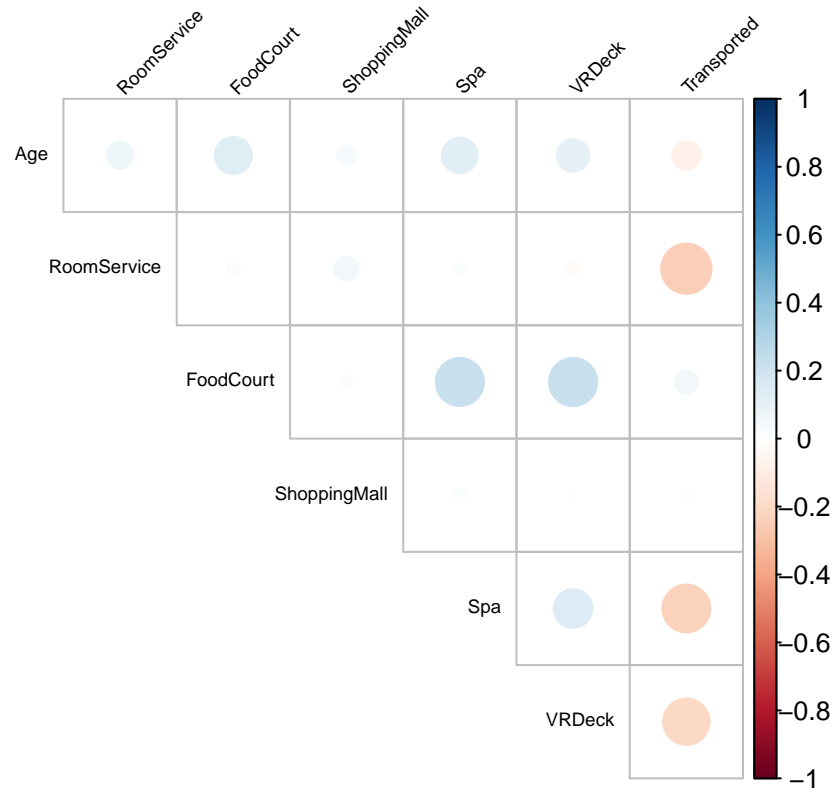
Because we have no reason to think that outliers are errors in data entry or significant data anomalies, and because our algorithms are relatively resistant to outliers, we do not remove outliers from the dataset.

C. Multicollinearity

While we were aware of the correlations with Transported, it is interesting to note that the correlations among various forms of spending are actually quite mild. We do not need to address multicollinearity in any systematic way.

```
dfNum <- dfTrain %>%  
  dplyr::select(where(is.numeric))  
dfNum <- na.omit(dfNum)  
  
a <- EHExplore_Multicollinearity(dfNum)
```


Heatmap for Multicollinearity Analysis

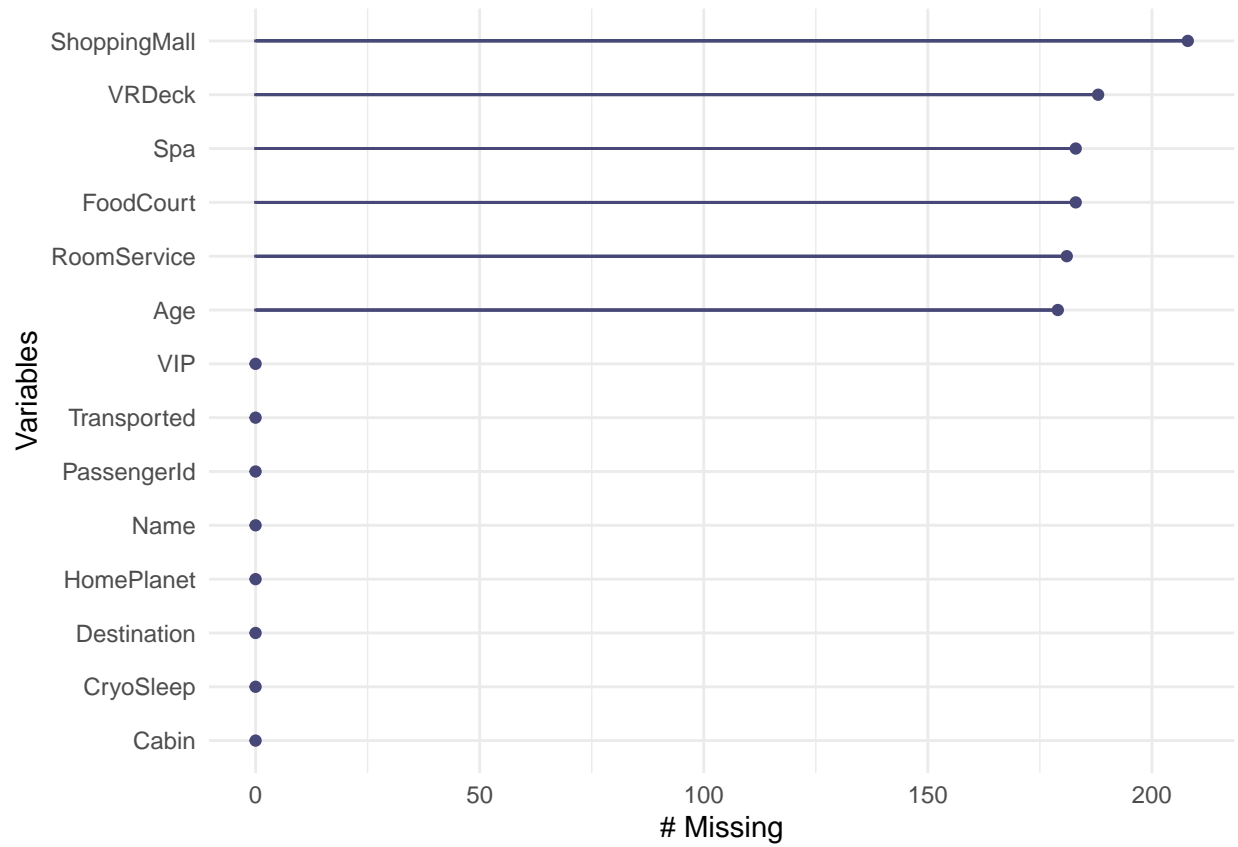


D. Missing Values

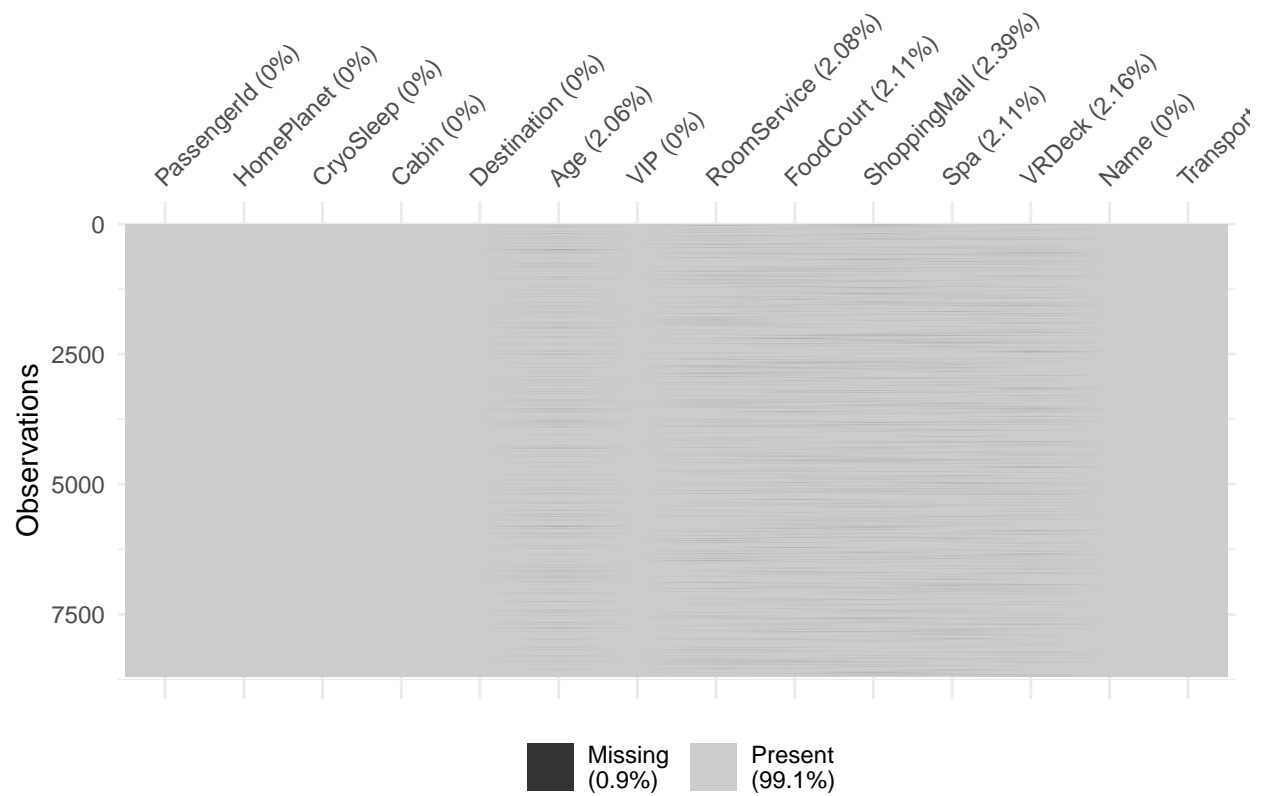
Missing values mainly appear for the amenities spending variables in the dataset. There are over 1,000 (12% of the database).

```
EHSummarize_MissingValues(dfTrain)
```

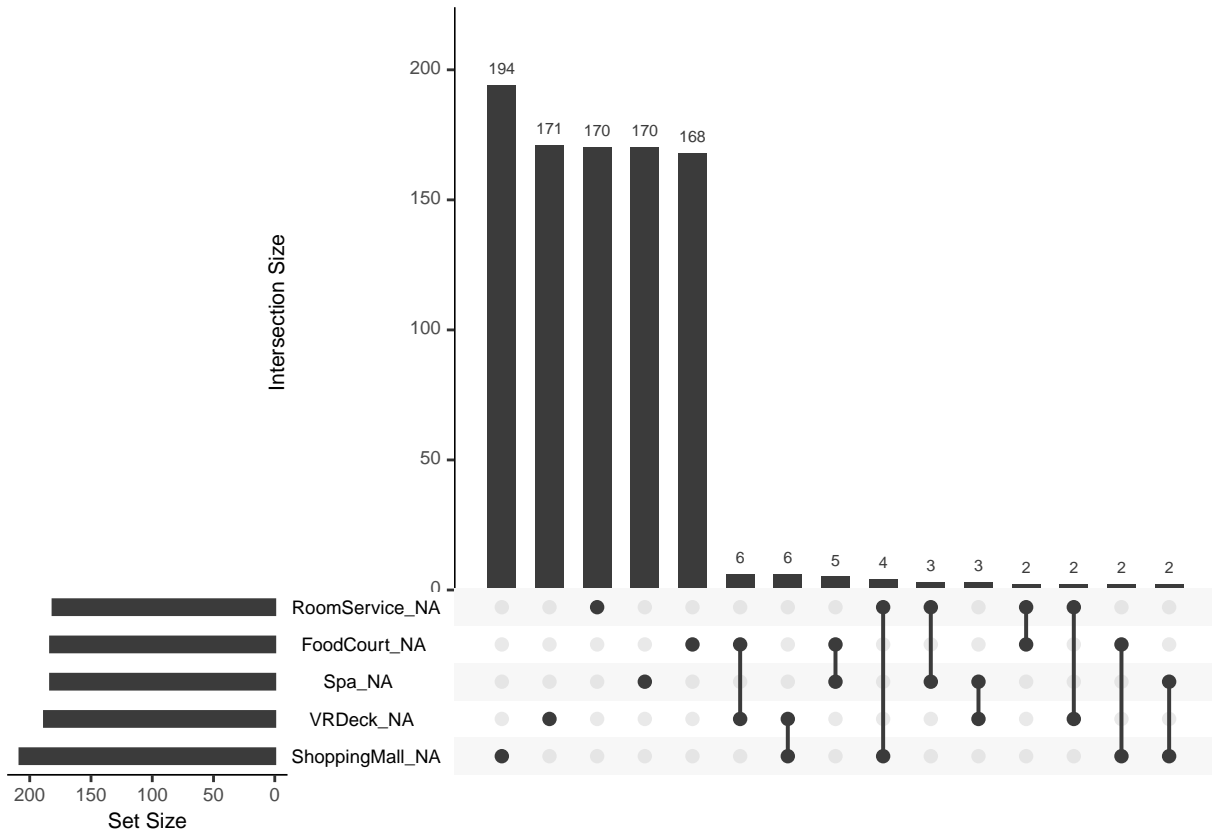
```
## [[1]]
```



```
##  
## [[2]]
```



```
##
## [[3]]
```



The missing values are not correlated with each other, suggesting they are probably missing at random. To further support this hypothesis we create flags for missing values and perform Chi Square tests against the target variable. None of the flags are significant. We will therefore remove the records with missing values from the training set (we will do this after some feature engineering), and impute the median for the test set.

```
dfMissingRecordsFlagAny <- dfTrain %>%
  mutate(Flag=ifelse(rowSums(is.na(dfTrain)) > 0, 1, 0)) %>%
  dplyr::select(Transported, Flag)

dfMissingRecordsFlag_RoomService <- dfTrain %>%
  mutate(RoomServiceFlag=ifelse(is.na(dfTrain$RoomService) > 0, 1, 0)) %>%
  dplyr::select(Transported, RoomServiceFlag)

dfMissingRecordsFlag_VRDeck <- dfTrain %>%
  mutate(VRDeckFlag=ifelse(is.na(dfTrain$VRDeck) > 0, 1, 0)) %>%
  dplyr::select(Transported, VRDeckFlag)

dfMissingRecordsFlag_SPA <- dfTrain %>%
  mutate(SpaFlag=ifelse(is.na(dfTrain$Spa) > 0, 1, 0)) %>%
  dplyr::select(Transported, SpaFlag)

dfMissingRecordsFlag_ShoppingMall <- dfTrain %>%
```

```

mutate(ShoppingMallFlag=ifelse(is.na(dfTrain$ShoppingMall) > 0, 1, 0)) %>%
  dplyr::select(Transported, ShoppingMallFlag)

dfMissingRecordsFlag_FoodCourt <- dfTrain %>%
  mutate(FoodCourtFlag=ifelse(is.na(dfTrain$FoodCourt) > 0, 1, 0)) %>%
  dplyr::select(Transported, FoodCourtFlag)

print(chisq.test(table(dfMissingRecordsFlagAny)))

```

```

##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  table(dfMissingRecordsFlagAny)
## X-squared = 0.15887, df = 1, p-value = 0.6902

```

```

print(chisq.test(table(dfMissingRecordsFlag_SPA)))

```

```

##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  table(dfMissingRecordsFlag_SPA)
## X-squared = 0.0098187, df = 1, p-value = 0.9211

```

```

print(chisq.test(table(dfMissingRecordsFlag_FoodCourt)))

```

```

##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  table(dfMissingRecordsFlag_FoodCourt)
## X-squared = 0.89665, df = 1, p-value = 0.3437

```

```

print(chisq.test(table(dfMissingRecordsFlag_VRDeck)))

```

```

##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  table(dfMissingRecordsFlag_VRDeck)
## X-squared = 0.1728, df = 1, p-value = 0.6776

```

```

print(chisq.test(table(dfMissingRecordsFlag_ShoppingMall)))

```

```

##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  table(dfMissingRecordsFlag_ShoppingMall)
## X-squared = 1.5072, df = 1, p-value = 0.2196

```

```
print(chisq.test(table(dfMissingRecordsFlag_RoomService)))
```

```
##
## Pearson's Chi-squared test with Yates' continuity correction
##
## data:  table(dfMissingRecordsFlag_RoomService)
## X-squared = 1.3229, df = 1, p-value = 0.2501
```

E. First Pass Logistic Regression: 9th percentile

We perform a logistic regression with what we have and post to Kaggle just to get a baseline. Accuracy on training is 77%, significantly better than the 51% no information rate, but gives us only 69% on the Kaggle set (which puts us at the 9th percentile).

```
dfTrainNum <- dfTrain %>%
  dplyr::select(where(is.numeric))

dfTestNum <- dfTest %>%
  dplyr::select(where(is.numeric))

dfTestNum <- cbind(dfTestNum, "PassengerID" = dfTest$PassengerId)

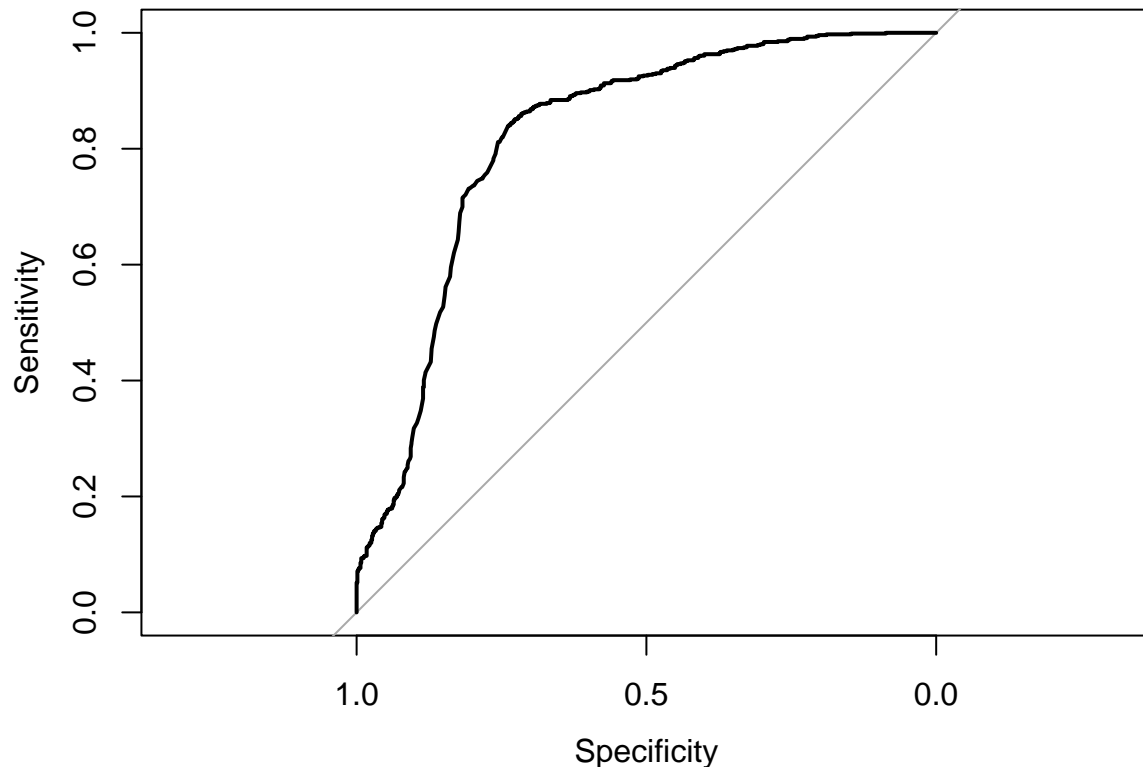
q <- EHModel_Regression_Logistic(dfTrainNum, "Transported", returnLM = TRUE)
```

```
##
## Call:
## glm(formula = fla, family = "binomial", data = train_reg)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4474  -0.8358   0.0166   0.8739   4.9443
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  7.125e-01  6.607e-02  10.785 < 2e-16 ***
## Age          2.085e-03  2.123e-03   0.982  0.32593
## RoomService -2.216e-03  1.108e-04 -20.008 < 2e-16 ***
## FoodCourt    7.701e-04  4.622e-05  16.662 < 2e-16 ***
## ShoppingMall 1.940e-04  6.630e-05   2.926  0.00343 **
## Spa         -2.488e-03  1.342e-04 -18.536 < 2e-16 ***
## VRDeck      -2.078e-03  1.165e-04 -17.837 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 8471.5  on 6110  degrees of freedom
## Residual deviance: 6288.6  on 6104  degrees of freedom
## (844 observations deleted due to missingness)
## AIC: 6302.6
##
## Number of Fisher Scoring iterations: 7
```

```

##
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 499  89
##           1 251 670
##
##           Accuracy : 0.7747
##           95% CI : (0.7528, 0.7955)
##           No Information Rate : 0.503
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.5488
##
## Mcnemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.6653
##           Specificity : 0.8827
##           Pos Pred Value : 0.8486
##           Neg Pred Value : 0.7275
##           Prevalence : 0.4970
##           Detection Rate : 0.3307
##           Detection Prevalence : 0.3897
##           Balanced Accuracy : 0.7740
##
##           'Positive' Class : 0
##

```



```
##
## Call:
## roc.default(response = dfPred_raw$class, predictor = dfPred_raw$predict_reg, plot = TRUE)
##
## Data: dfPred_raw$predict_reg in 750 controls (dfPred_raw$class 0) < 759 cases (dfPred_raw$class 1).
## Area under the curve: 0.822

predictions_logistic <- EHModel_Predict(q, dfTestNum, threshold=.5, testData_IDColumn = "PassengerID",
predictions_logistic$Transported <- ifelse(predictions_logistic$Transported==1,"True","False")
write_csv(predictions_logistic, "C://Users//erico//Desktop//LR_FirstPass.csv")
```

Data Preparation and Feature Engineering

1. Create groups based on the Passenger ID

Passenger IDs are constructed to identify passengers travelling in groups. We create groupings from the ID. How likely is it that if the majority of members of a group were transported, then they all were transported? Only somewhat likely. A histogram shows the distribution of percentages of transported within groups. Most often, half the members transported and half did not.

```
Group2 <- dfTrain %>% dplyr::select(PassengerId, Transported)
```



```

Group2$Group <- sub("\\_.*", "", Group2$PassengerId)

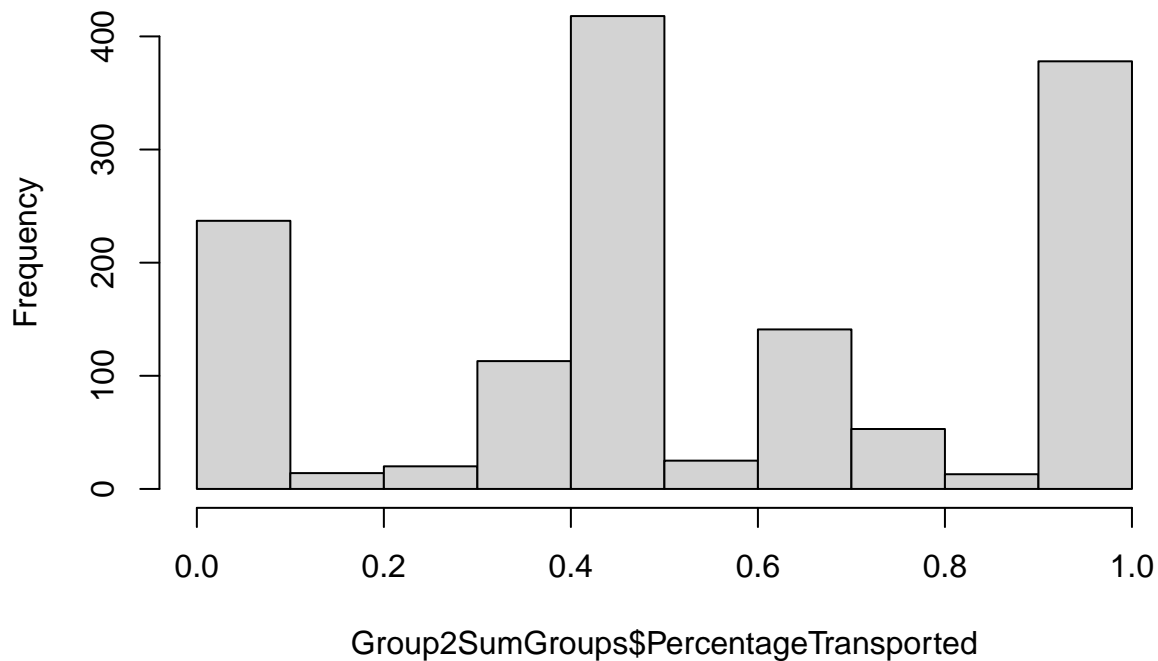
Group2Sum <- Group2 %>%
  dplyr::group_by(Group) %>%
  dplyr::summarize(PercentageTransported= sum(Transported)/dplyr::n(), count=dplyr::n())

Group2SumGroups <- Group2Sum %>%
  filter(count>1)

hist(Group2SumGroups$PercentageTransported)

```

Histogram of Group2SumGroups\$PercentageTransported



```

dfTrainA <- dfTrain

dfTrainA$Group <- Group2$Group

dfTrainAA <- inner_join(dfTrainA, Group2Sum, by="Group") %>% dplyr::select(-Group, -PercentageTransported)

dfTrainBB <- dfTrainAA

#boxplot(as.numeric(dfTrainBB$count), as.numeric(dfTrainBB$Transported))

Group2a <- dfTest %>% dplyr::select(PassengerId)

Group2a$Group <- sub("\\_.*", "", Group2a$PassengerId)

```

```

Group2aSum <- Group2a %>%
  dplyr::group_by(Group) %>%
  dplyr::summarize(count=dplyr::n())

dfTestA <- dfTest

dfTestA$Group <- Group2a$Group

dfTestAA <- inner_join(dfTestA, Group2aSum, by="Group") %>% dplyr::select(-Group)

dfTrainAA$InAGroup = ifelse(dfTrainAA$count>1,1,0)
dfTestAA$InAGroup = ifelse(dfTestAA$count>1,1,0)

dfTrain1 <- na.omit(dfTrainAA)
dfTest1 <- EHPprepare_MissingValues_Imputation(dfTestAA, impute="median")

```

2. Create Cabin Variables

Cabin variables consist of 3 parts in the form of a/b/c which indicate the location of the cabin on the ship. Here we extract out parts “a” and “c” - b appears to have no influence on the target.

```

library(stringr)
dfTrain1$Cabin1 <- substr(dfTrain1$Cabin,1,1)
dfTrain1$Cabin2 <- str_sub(dfTrain1$Cabin, - 1, - 1)

dfTrain2 <- dfTrain1 %>%
  dplyr::select(-Cabin)

dfTest1$Cabin1 <- substr(dfTest1$Cabin,1,1)
dfTest1$Cabin2 <- str_sub(dfTest1$Cabin, - 1, - 1)

dfTest2 <- dfTest1 %>%
  dplyr::select(-Cabin)

dfTrain3 <- dfTrain2
dfTest3 <- dfTest2

```

3. Create Dummy Variables

Now that we have engineered Cabin, we create dummy variables to handle category variables throughout the dataset.

```

dfTrain4 <- dfTrain3 %>%
  dplyr::select(-Name, -PassengerId)

PassengerID <- dfTest3$PassengerId
dfTest4 <- dfTest3 %>%
  dplyr::select(-Name, -PassengerId)

dfTrain5v <- EHPprepare_CreateDummies(dfTrain4, "Transported")
dfTrain5 <- dfTrain5v %>%
  dplyr::select(-Cabin2_)

```

```
dfTest5v <- EHPPrepare_CreateDummies(dfTest4, "Transported")
dfTest5 <- dfTest5v %>%
  dplyr::select(-Cabin2_)

dfTest5 <- cbind(PassengerID, dfTest5)
```

4. Implement Interaction Features

There were a number of interactions among variables which appeared when the variables were examined in isolation. We only retained one (group passengers were more likely to be transported if they shopped at the mall) because the others did not significantly increase accuracy when running the overall model - but there would be more to explore here.

```
dfTrain5$Inter_CountShop = dfTrain5$InAGroup*dfTrain5$ShoppingMall
dfTest5$Inter_CountShop = dfTest5$InAGroup*dfTest5$ShoppingMall
```

5. Perform Logistic Regression With Engineered Features: 26th Percentile

We perform a logistic regression with what we have and post to Kaggle just to get a baseline. Accuracy on the holdout set improves training to 79% (compared to 77% on the untransformed regression), and gives us 78% on the Kaggle set, which puts us at the 26th percentile.

```
dfTrain5z <- dfTrain5

dfTrain5zz <- dfTrain5z %>%
  dplyr::select(-count)

dfTest5z <- dfTest5

dfTest5zz <- dfTest5z %>%
  dplyr::select(-count)

set.seed(042758)

q <- EHModel_Regression_Logistic(dfTrain5zz, "Transported", returnLM = TRUE)
```

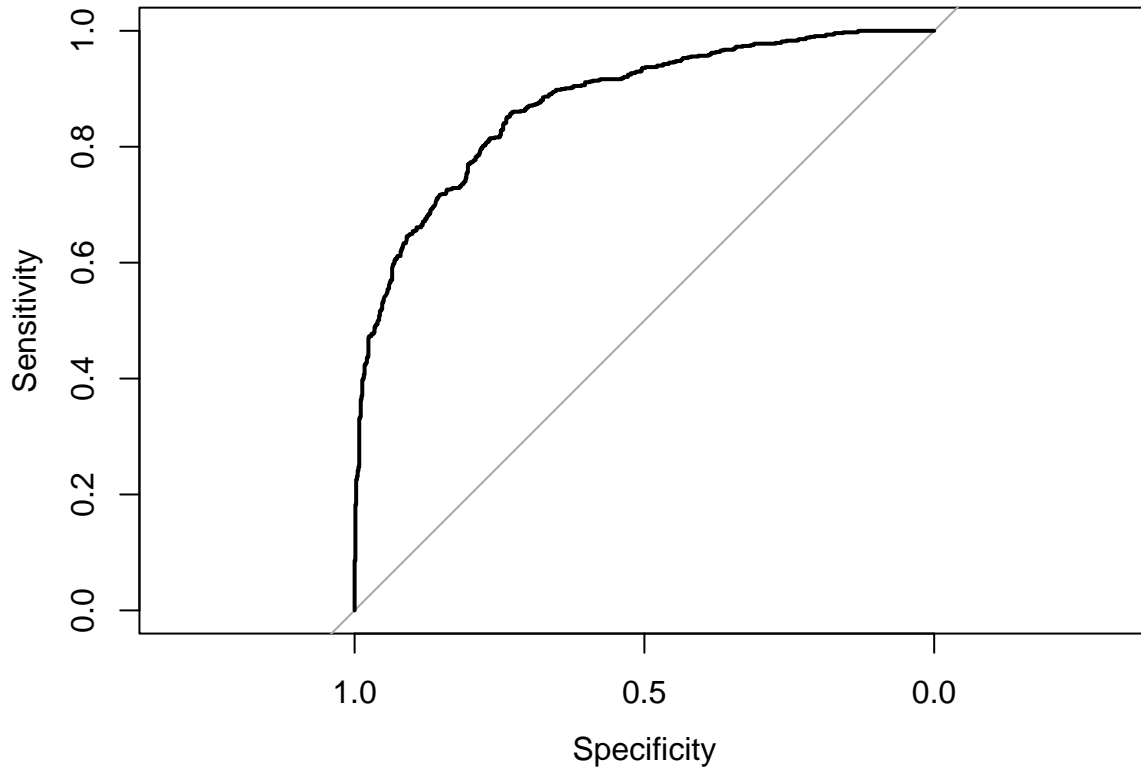
```
##
## Call:
## glm(formula = fla, family = "binomial", data = train_reg)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.9252  -0.6568   0.0149   0.6997   3.3536
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.2955844  0.1180833   2.503 0.012308 *
## Age           -0.0079509  0.0025348  -3.137 0.001709 **
## RoomService   -0.0015488  0.0001113 -13.920 < 2e-16 ***
## FoodCourt      0.0006140  0.0000530  11.585 < 2e-16 ***
## ShoppingMall    0.0006873  0.0001026   6.701 2.07e-11 ***
```

```

## Spa -0.0020068 0.0001222 -16.427 < 2e-16 ***
## VRDeck -0.0019359 0.0001217 -15.910 < 2e-16 ***
## InAGroup 0.1753772 0.0802973 2.184 0.028955 *
## HomePlanet_ 0.4533562 0.2195945 2.065 0.038969 *
## HomePlanet_Europa 1.7502630 0.2681931 6.526 6.75e-11 ***
## HomePlanet_Mars 0.5704315 0.1141614 4.997 5.83e-07 ***
## CryoSleep_ 0.2728606 0.2061154 1.324 0.185561
## CryoSleep_True 1.3798768 0.0978384 14.104 < 2e-16 ***
## Destination_ 0.4335945 0.2398617 1.808 0.070655 .
## Destination_55.Cancric.e 0.4794132 0.0980725 4.888 1.02e-06 ***
## Destination_PSO.J318.5.22 0.0326575 0.1102523 0.296 0.767072
## VIP_ 0.2137313 0.2274328 0.940 0.347343
## VIP_True -0.1144653 0.3008491 -0.380 0.703593
## Cabin1_ -0.4472731 0.2413786 -1.853 0.063883 .
## Cabin1_A -0.9518680 0.3502360 -2.718 0.006572 **
## Cabin1_B 0.2368897 0.3155204 0.751 0.452779
## Cabin1_C 1.3388313 0.3520418 3.803 0.000143 ***
## Cabin1_D -0.1615527 0.2077058 -0.778 0.436689
## Cabin1_E -0.6671090 0.1240673 -5.377 7.57e-08 ***
## Cabin1_G -0.4013448 0.0994084 -4.037 5.41e-05 ***
## Cabin1_T -1.3950991 1.8765514 -0.743 0.457217
## Cabin2_P -0.6597095 0.0704040 -9.370 < 2e-16 ***
## Inter_CountShop -0.0003796 0.0001440 -2.637 0.008368 **
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 8450.7 on 6095 degrees of freedom
## Residual deviance: 5172.8 on 6068 degrees of freedom
## AIC: 5228.8
##
## Number of Fisher Scoring iterations: 7
##
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 580 143
##           1 177 624
##
##           Accuracy : 0.79
##           95% CI : (0.7687, 0.8102)
##           No Information Rate : 0.5033
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.5799
##
## Mcnemar's Test P-Value : 0.06507
##
##           Sensitivity : 0.7662
##           Specificity : 0.8136
##           Pos Pred Value : 0.8022
##           Neg Pred Value : 0.7790

```

```
##           Prevalence : 0.4967
##           Detection Rate : 0.3806
##           Detection Prevalence : 0.4744
##           Balanced Accuracy : 0.7899
##
##           'Positive' Class : 0
##
```



```
##
## Call:
## roc.default(response = dfPred_raw$class, predictor = dfPred_raw$predict_reg, plot = TRUE)
##
## Data: dfPred_raw$predict_reg in 757 controls (dfPred_raw$class 0) < 767 cases (dfPred_raw$class 1).
## Area under the curve: 0.8751
```

```
predictions_logistic <- EHModel_Predict(q, dfTest5zz, threshold=.5, testData_IDColumn = "PassengerID",
predictions_logistic$Transported <- ifelse(predictions_logistic$Transported==1,"True","False")
write_csv(predictions_logistic, "C://Users//erico//Desktop//LR_WithFeatures.csv")
```

At this point, a picture of the transported passengers begins to emerge - they are the poorer passengers, most likely to shop on a budget or, even cheaper, spend the trip in cryosleep. They tend to enter the ship in groups and inhabit lower class cabins.

More Complex Models

Given the apparent complexity of the data shape, we turn to more complex nonparametric models to improve our predictions.

1. Perform SVM: 70th Percentile

We begin with Support Vector Machines and try three kernels - linear, poly and radial. We perform ten-fold cross validation and optimal hypertuning of C based on the caret package. Radial performs the best (accuracy=80.1% on the kaggle set) and boosts us to the 70th percentile.

```
library(doParallel)
library(caret)

cl<-makePSOCKcluster(7)

registerDoParallel(cl)

print("Linear: -----")

## [1] "Linear: -----"

lin <- EHModel_SVM(dfTrain5zz, "Transported", method="linear")

## Support Vector Machines with Linear Kernel
##
## 6097 samples
## 27 predictor
## 2 classes: '0', '1'
##
## Pre-processing: centered (27), scaled (27)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 5487, 5487, 5487, 5488, 5488, ...
## Resampling results across tuning parameters:
##
##      C          Accuracy  Kappa
## 0.0100000 0.7760621 0.5524086
## 0.1147368 0.7987524 0.5974004
## 0.2194737 0.8002288 0.6003060
## 0.3242105 0.8000653 0.5999630
## 0.4289474 0.8004484 0.6007181
## 0.5336842 0.8001749 0.6001643
## 0.6384211 0.8001748 0.6001610
## 0.7431579 0.8002845 0.6003763
## 0.8478947 0.8003938 0.6005952
## 0.9526316 0.8004481 0.6007033
## 1.0573684 0.8001202 0.6000454
## 1.1621053 0.8004484 0.6007013
## 1.2668421 0.8002297 0.6002636
## 1.3715789 0.8002845 0.6003721
## 1.4763158 0.8002845 0.6003719
## 1.5810526 0.8002844 0.6003714
```

```

## 1.6857895 0.8002845 0.6003698
## 1.7905263 0.8005032 0.6008077
## 1.8952632 0.8003390 0.6004790
## 2.0000000 0.8003937 0.6005881
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was C = 1.790526.
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 562 133
##           1 195 633
##
##           Accuracy : 0.7846
##           95% CI : (0.7631, 0.805)
##           No Information Rate : 0.503
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.569
##
## Mcnemar's Test P-Value : 0.0007567
##
##           Sensitivity : 0.7424
##           Specificity : 0.8264
##           Pos Pred Value : 0.8086
##           Neg Pred Value : 0.7645
##           Prevalence : 0.4970
##           Detection Rate : 0.3690
##           Detection Prevalence : 0.4563
##           Balanced Accuracy : 0.7844
##
##           'Positive' Class : 0
##

```

```
print("Poly: -----")
```

```
## [1] "Poly: -----"
```

```
poly <- EHModel_SVM(dfTrain5zz, "Transported", method="poly")
```

```

## Support Vector Machines with Polynomial Kernel
##
## 6097 samples
## 27 predictor
## 2 classes: '0', '1'
##
## Pre-processing: centered (27), scaled (27)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 5487, 5487, 5487, 5488, 5488, 5488, ...
## Resampling results across tuning parameters:
##
## degree scale C      Accuracy  Kappa

```

```

## 1      0.001 0.25 0.7551225 0.5108272
## 1      0.001 0.50 0.7498187 0.5002867
## 1      0.001 1.00 0.7554498 0.5114740
## 1      0.010 0.25 0.7633219 0.5271118
## 1      0.010 0.50 0.7703213 0.5410180
## 1      0.010 1.00 0.7759529 0.5521904
## 1      0.100 0.25 0.7869427 0.5740405
## 1      0.100 0.50 0.7943245 0.5886851
## 1      0.100 1.00 0.7984247 0.5967629
## 2      0.001 0.25 0.7497093 0.5000669
## 2      0.001 0.50 0.7552857 0.5111456
## 2      0.001 1.00 0.7615722 0.5236397
## 2      0.010 0.25 0.7717418 0.5438533
## 2      0.010 0.50 0.7814206 0.5631022
## 2      0.010 1.00 0.7927376 0.5855954
## 2      0.100 0.25 0.8023051 0.6045642
## 2      0.100 0.50 0.8050391 0.6099812
## 2      0.100 1.00 0.8044376 0.6087612
## 3      0.001 0.25 0.7527161 0.5060429
## 3      0.001 0.50 0.7588931 0.5183145
## 3      0.001 1.00 0.7650716 0.5305831
## 3      0.010 0.25 0.7802723 0.5608270
## 3      0.010 0.50 0.7912077 0.5825486
## 3      0.010 1.00 0.7976053 0.5952331
## 3      0.100 0.25 0.7935057 0.5869583
## 3      0.100 0.50 0.7950370 0.5900214
## 3      0.100 1.00 0.7952015 0.5903619
##
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were degree = 2, scale = 0.1 and C = 0.5.
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 570 148
##           1 187 618
##
##           Accuracy : 0.78
##           95% CI : (0.7584, 0.8006)
##           No Information Rate : 0.503
##           P-Value [Acc > NIR] : < 2e-16
##
##           Kappa : 0.5599
##
## McNemar's Test P-Value : 0.03788
##
##           Sensitivity : 0.7530
##           Specificity : 0.8068
##           Pos Pred Value : 0.7939
##           Neg Pred Value : 0.7677
##           Prevalence : 0.4970
##           Detection Rate : 0.3743
##           Detection Prevalence : 0.4714
##           Balanced Accuracy : 0.7799

```



```

##
##      'Positive' Class : 0
##

print("Radial: -----")

## [1] "Radial: -----"

rad <- EHModel_SVM(dfTrain5zz, "Transported", method="radial")

## Support Vector Machines with Radial Basis Function Kernel
##
## 6097 samples
## 27 predictor
## 2 classes: '0', '1'
##
## Pre-processing: centered (27), scaled (27)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 5487, 5487, 5487, 5488, 5488, 5488, ...
## Resampling results across tuning parameters:
##
##  C      Accuracy  Kappa
##  0.25  0.7841545  0.5685024
##  0.50  0.7899499  0.5799809
##  1.00  0.7951430  0.5902743
##
## Tuning parameter 'sigma' was held constant at a value of 0.04367297
## Accuracy was used to select the optimal model using the largest value.
## The final values used for the model were sigma = 0.04367297 and C = 1.
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 574 165
##           1 183 601
##
##           Accuracy : 0.7715
##           95% CI : (0.7496, 0.7924)
##           No Information Rate : 0.503
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.5429
##
## Mcnemar's Test P-Value : 0.3621
##
##           Sensitivity : 0.7583
##           Specificity : 0.7846
##           Pos Pred Value : 0.7767
##           Neg Pred Value : 0.7666
##           Prevalence : 0.4970
##           Detection Rate : 0.3769
##           Detection Prevalence : 0.4852
##           Balanced Accuracy : 0.7714

```

```
##
##      'Positive' Class : 0
##

predictions_rad <- EHModel_Predict(rad$svm, dfTest5zz, predictionsColumnName = "Transported", testData = dfTest5zz)

predictions_rad$Transported <- ifelse(predictions_rad$Transported==1,"True","False")
write_csv(predictions_rad, "C://Users//erico//Desktop//SVM_Rad.csv")

stopCluster(cl)
```

2. Perform limited Neural Networks

Neural Networks require a great deal of computer resources and time. A simple first pass with two hidden layers took a great deal of time, needed a high stepmax to converge and provided poor results (15th percentile). The algorithm was therefore difficult to hypertune.

```
library(neuralnet)
set.seed(42760)

dfTrainN <- dfTrain5zz
dfTrainN$Transported = as.factor(dfTrainN$Transported)
dfScale <- EHPrepare_ScaleAllButTarget(dfTrainN, "Transported")

#n <- neuralnet(Transported ~ .,
  #data = dfScale,
  # hidden = 2,
  #err.fct = "ce",
  # linear.output = F,
  # lifesign = 'full',
  # rep = 2,
  # algorithm = "rprop+",
  # stepmax = 100000, likelihood=TRUE, )

#plot(n, rep = 2)

#predicts <- predict(n, dfTest5zz)
#predict_nn <- as.data.frame(predict(n, dfTest5zz)) %>%
#  dplyr::select(V1) %>%
#  dplyr::rename("Transported" = V1) %>%
#  mutate(Transported=as.numeric(Transported)) %>%
#  mutate(Transported=ifelse(Transported>.5,"True","False"))

#predictions_nn <- EHModel_Predict(n, dfTest5zz, predictionsColumnName = "Transported", #testData_IDCol = dfTest5zz$IDCol)

#write_csv(predictions_nn, "C://Users//erico//Desktop//nn.csv")
```

In order to address the long time until convergence (many hours), we experimented with dimensionality reduction, but this was ineffective. PCA, e.g., did not result in a small number of components taking the largest share of variance. Taking a sample of records or manually eliminating columns allowed for faster run times but hurt performance. Below is the result of PCA analysis:

```
pca <- prcomp(dfTrain5zz, center = TRUE, scale. = TRUE)

summary(pca)
```

```
## Importance of components:
##
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
## Standard deviation	1.7113	1.50279	1.32219	1.24229	1.08522	1.07691	1.06709
## Proportion of Variance	0.1046	0.08066	0.06243	0.05512	0.04206	0.04142	0.04067
## Cumulative Proportion	0.1046	0.18524	0.24768	0.30279	0.34486	0.38627	0.42694

```
##
```

	PC8	PC9	PC10	PC11	PC12	PC13	PC14
## Standard deviation	1.05502	1.0382	1.02008	1.01195	1.00481	0.99789	0.98678
## Proportion of Variance	0.03975	0.0385	0.03716	0.03657	0.03606	0.03556	0.03478
## Cumulative Proportion	0.46669	0.5052	0.54236	0.57893	0.61499	0.65055	0.68533

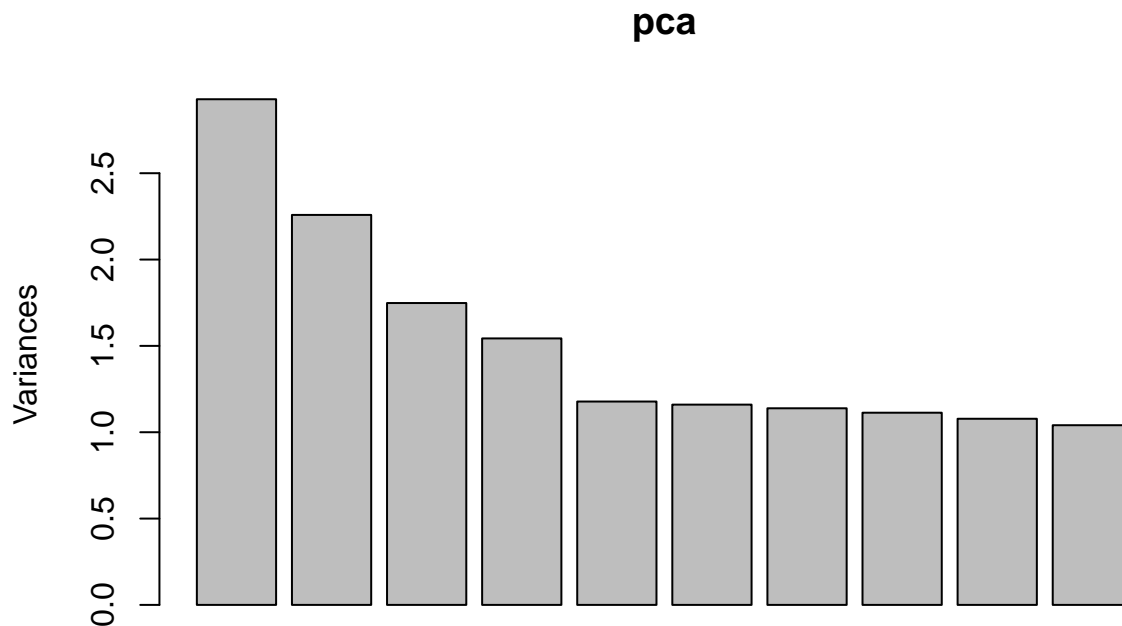
```
##
```

	PC15	PC16	PC17	PC18	PC19	PC20	PC21
## Standard deviation	0.98295	0.93825	0.93370	0.9150	0.90835	0.88785	0.87515
## Proportion of Variance	0.03451	0.03144	0.03114	0.0299	0.02947	0.02815	0.02735
## Cumulative Proportion	0.71983	0.75127	0.78241	0.8123	0.84178	0.86993	0.89729

```
##
```

	PC22	PC23	PC24	PC25	PC26	PC27	PC28
## Standard deviation	0.81848	0.8111	0.76818	0.64984	0.50335	0.45711	0.27096
## Proportion of Variance	0.02393	0.0235	0.02108	0.01508	0.00905	0.00746	0.00262
## Cumulative Proportion	0.92121	0.9447	0.96578	0.98087	0.98992	0.99738	1.00000

```
plot(pca)
```



3. Tree Algorithms 1: Perform Random Forest: 34th Percentile

Random forest is our first tree. We use the parallel library to run the model on multiple cores.

We perform ten-fold cross validation and optimal hypertuning of sigma and c based on the caret package.

This gives us 79% on the holdout set and 78.8% on the kaggle set, which puts us in the 34th percentile. This is low compared to SVM.

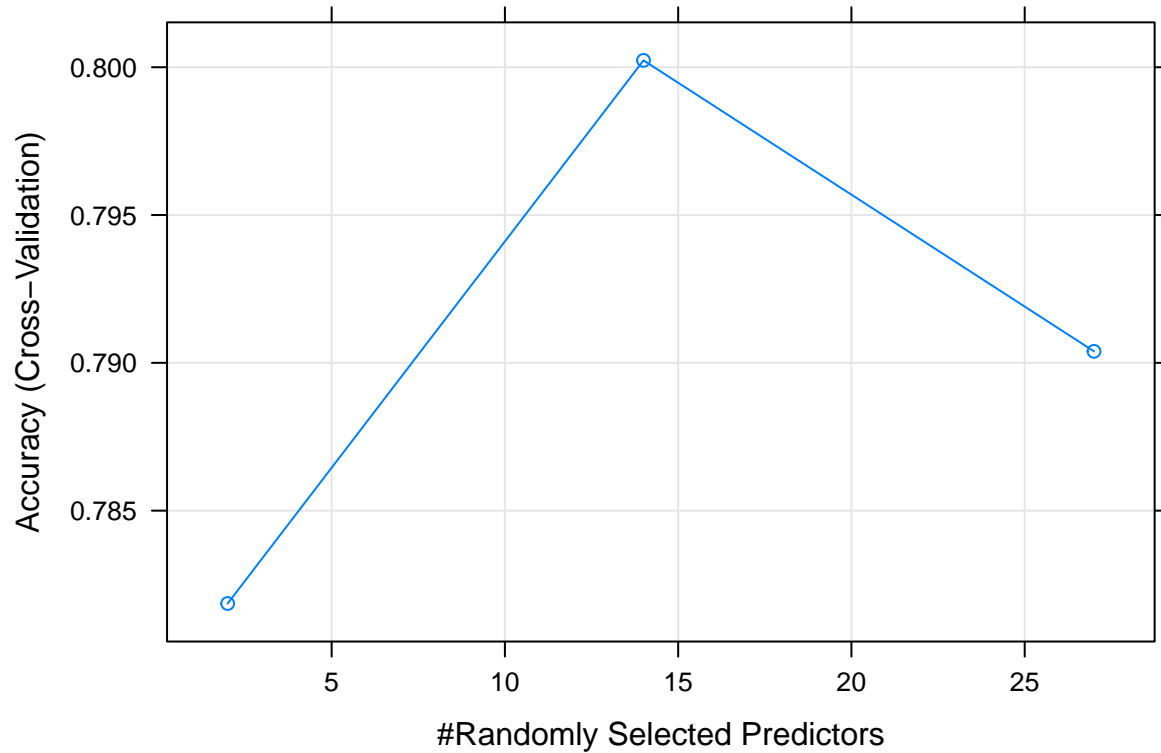
```
library(doParallel)

cl<-makePSOCKcluster(7)

registerDoParallel(cl)

a <- EHModel_RandomForest(dfTrain5zz, "Transported")

## Random Forest
##
## 6097 samples
## 27 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 5487, 5487, 5487, 5488, 5488, 5488, ...
## Resampling results across tuning parameters:
##
##  mtry  Accuracy  Kappa
##    2    0.7818587 0.5639198
##   14    0.8002320 0.6005062
##   27    0.7903895 0.5809306
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 14.
```



```
## rf variable importance
##
##   only 20 most important variables shown (out of 27)
##
##               Overall
## CryoSleep_True 100.000
## Spa             85.968
## Age             85.931
## VRDeck          80.377
## RoomService     76.410
## FoodCourt       67.359
## ShoppingMall    45.336
## Cabin1_G        20.058
## Cabin2_P        15.122
## Cabin1_E        13.814
## HomePlanet_Europa 12.745
## Inter_CountShop 10.717
## Destination_55.Cancrile 9.720
## InAGroup        9.616
## HomePlanet_Mars  8.855
## Destination_PS0.J318.5.22 8.058
## HomePlanet_     4.233
## CryoSleep_      3.845
## VIP_            3.684
## Cabin1_C        3.332
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   0   1
##           0 606 169
##           1 151 597
##
##           Accuracy : 0.7899
##           95% CI : (0.7686, 0.8101)
##           No Information Rate : 0.503
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.5798
##
## Mcnemar's Test P-Value : 0.3419
##
##           Sensitivity : 0.8005
##           Specificity : 0.7794
##           Pos Pred Value : 0.7819
##           Neg Pred Value : 0.7981
##           Prevalence : 0.4970
##           Detection Rate : 0.3979
##           Detection Prevalence : 0.5089
##           Balanced Accuracy : 0.7900
##
##           'Positive' Class : 0
##
## [1] "Parameters:  mtry = 14 , ntree = 500 , nrnodes = 1605"
```

```
predictions_rf <- EHModel_Predict(a$rfr, dfTest5zz, predictionsColumnName = "Transported", testData_IDCol = "ID")

predictions_rf$Transported <- ifelse(predictions_rf$Transported==1,"True","False")
write_csv(predictions_rf, "C://Users//erico//Desktop//RF1.csv")

stopCluster(cl)
```

4. Tree Algorithms 2: Perform XGBoost Untuned: 73rd Percentile

At the 34th percentile we need a more powerful model. As an active learner, XGBoost is likely to fit our training model better than random forest, though it may overfit the data.

Our untuned model, with 55 rounds (chosen randomly), achieves 79% accuracy on the holdout set and 80.2% accuracy on the kaggle set, and reaches the 74th percentile.

```
library(xgboost)
library(caret)

set.seed(0)

dfTrain5q <- dfTrain5
dfTest5q <- dfTest5

dfTrain5q$Transported = as.character(dfTrain5q$Transported)
```

```

dfTrain5$Transported = as.character(dfTrain5$Transported)

train_x <- dfTrain5q %>%
  dplyr::select(-Transported, -count)
train_y <- dfTrain5q %>%
  dplyr::select(Transported)

dfTest6 <- dfTest5q
dfTest6$Transported = "1"

test_x <- dfTest6 %>%
  dplyr::select(-Transported, -PassengerID, -count)
test_y <- dfTest6 %>%
  dplyr::select(Transported)

xgb_train <- xgb.DMatrix(data = as.matrix(train_x), label = as.matrix(train_y))
xgb_test <- xgb.DMatrix(data = as.matrix(test_x), label = as.matrix(test_y))

#depth of 1000 and nrounds 100000 showed no double descent

model2 <- xgb.train(data = xgb_train, max.depth = 3, nrounds = 50)

predictions <- predict(model2, newdata=xgb_test)
predictions <- data.frame(as.vector(predictions))
predictions$PassengerId <- dfTest5$PassengerID
predictions[,c(1,2)] <- predictions[,c(2,1)]
colnames(predictions) <- c("PassengerID", "Transported")

predictions <- predictions %>%
  mutate(Transported=ifelse(Transported>.5, 'True', 'False'))

write_csv(predictions, "C:\\Users\\erico\\Desktop\\XGBpredictions_Rounds50.csv")

```

Hypertuning

Now we try some experiments with hypertuning. First we choose a more optimal nrounds for xgboost. Then we try manually adjusting sigma both up and down on our radial SVM model to see if that improves performance.

1. Tune XGBoost: 78th Percentile

We perform ten-fold cross validation on 1-100 nrounds and find the optimal nrounds for accuracy (14). Our model, with 14 rounds, achieves 80.3% accuracy and reaches the 74th percentile.

```

set.seed(100)
cv <- xgb.cv(data = xgb_train, nrounds = 100, nthread = 2, nfold = 10, metrics = list("rmse", "auc", "error"),
             max_depth = 3, eta = 1, objective = "binary:logistic")

```

```
## [1] train-rmse:0.415047+0.000841 train-auc:0.813023+0.001987 train-error:0.282094+0.001992 tes
```

## [2]	train-rmse:0.385845+0.001185	train-auc:0.860516+0.002116	train-error:0.222164+0.007927	tes
## [3]	train-rmse:0.378852+0.001398	train-auc:0.870728+0.001990	train-error:0.209274+0.003398	tes
## [4]	train-rmse:0.372729+0.001637	train-auc:0.880756+0.002273	train-error:0.203893+0.003009	tes
## [5]	train-rmse:0.368863+0.001861	train-auc:0.886316+0.002073	train-error:0.201575+0.002140	tes
## [6]	train-rmse:0.366441+0.001607	train-auc:0.889350+0.001927	train-error:0.199446+0.002347	tes
## [7]	train-rmse:0.364508+0.001586	train-auc:0.891987+0.001954	train-error:0.196544+0.003051	tes
## [8]	train-rmse:0.361741+0.002500	train-auc:0.895200+0.002572	train-error:0.192797+0.004358	tes
## [9]	train-rmse:0.359367+0.002674	train-auc:0.897881+0.002820	train-error:0.190580+0.005296	tes
## [10]	train-rmse:0.357505+0.002598	train-auc:0.900123+0.002623	train-error:0.188655+0.005272	tes
## [11]	train-rmse:0.356330+0.002516	train-auc:0.901504+0.002411	train-error:0.187474+0.005112	tes
## [12]	train-rmse:0.354930+0.002560	train-auc:0.903052+0.002540	train-error:0.185681+0.004593	tes
## [13]	train-rmse:0.353317+0.002560	train-auc:0.905031+0.002491	train-error:0.183129+0.003636	tes
## [14]	train-rmse:0.352609+0.002423	train-auc:0.905840+0.002382	train-error:0.181846+0.003486	tes
## [15]	train-rmse:0.351316+0.002204	train-auc:0.907134+0.002231	train-error:0.180782+0.003400	tes
## [16]	train-rmse:0.350333+0.002146	train-auc:0.908162+0.002161	train-error:0.179280+0.002695	tes
## [17]	train-rmse:0.349126+0.002166	train-auc:0.909447+0.002254	train-error:0.177924+0.003161	tes
## [18]	train-rmse:0.348295+0.002153	train-auc:0.910399+0.002297	train-error:0.177705+0.002517	tes
## [19]	train-rmse:0.347184+0.002109	train-auc:0.911515+0.002141	train-error:0.176363+0.002503	tes
## [20]	train-rmse:0.346005+0.001498	train-auc:0.912769+0.001518	train-error:0.175168+0.002447	tes
## [21]	train-rmse:0.345138+0.001769	train-auc:0.913644+0.001712	train-error:0.174161+0.003187	tes
## [22]	train-rmse:0.344329+0.001754	train-auc:0.914464+0.001689	train-error:0.173637+0.003224	tes
## [23]	train-rmse:0.343350+0.001915	train-auc:0.915459+0.001870	train-error:0.172907+0.003426	tes
## [24]	train-rmse:0.342475+0.001964	train-auc:0.916324+0.001901	train-error:0.171945+0.003266	tes
## [25]	train-rmse:0.341509+0.002072	train-auc:0.917216+0.002013	train-error:0.171289+0.003922	tes
## [26]	train-rmse:0.340689+0.001933	train-auc:0.918042+0.001888	train-error:0.170516+0.003508	tes
## [27]	train-rmse:0.339792+0.002083	train-auc:0.918946+0.002021	train-error:0.169714+0.004183	tes
## [28]	train-rmse:0.339044+0.002041	train-auc:0.919682+0.001946	train-error:0.169175+0.003845	tes
## [29]	train-rmse:0.338046+0.001954	train-auc:0.920647+0.001815	train-error:0.167760+0.003082	tes
## [30]	train-rmse:0.337128+0.001795	train-auc:0.921598+0.001656	train-error:0.166900+0.002514	tes
## [31]	train-rmse:0.336358+0.001755	train-auc:0.922325+0.001597	train-error:0.166229+0.002539	tes
## [32]	train-rmse:0.335645+0.001621	train-auc:0.922911+0.001497	train-error:0.165150+0.002449	tes
## [33]	train-rmse:0.334922+0.001529	train-auc:0.923596+0.001316	train-error:0.164071+0.002469	tes
## [34]	train-rmse:0.333932+0.001624	train-auc:0.924532+0.001408	train-error:0.163459+0.002936	tes
## [35]	train-rmse:0.332912+0.001554	train-auc:0.925443+0.001347	train-error:0.162321+0.002618	tes
## [36]	train-rmse:0.332262+0.001438	train-auc:0.926035+0.001194	train-error:0.161738+0.002974	tes
## [37]	train-rmse:0.331246+0.001364	train-auc:0.926959+0.001162	train-error:0.161534+0.002805	tes
## [38]	train-rmse:0.330345+0.001437	train-auc:0.927820+0.001221	train-error:0.159740+0.003067	tes
## [39]	train-rmse:0.329495+0.001612	train-auc:0.928612+0.001336	train-error:0.158953+0.003136	tes
## [40]	train-rmse:0.328807+0.001581	train-auc:0.929218+0.001320	train-error:0.158093+0.003108	tes
## [41]	train-rmse:0.327937+0.001694	train-auc:0.930023+0.001379	train-error:0.157480+0.003106	tes
## [42]	train-rmse:0.327184+0.001656	train-auc:0.930684+0.001377	train-error:0.156416+0.003253	tes
## [43]	train-rmse:0.326425+0.001521	train-auc:0.931303+0.001260	train-error:0.155497+0.002907	tes
## [44]	train-rmse:0.325603+0.001620	train-auc:0.931990+0.001330	train-error:0.154666+0.002618	tes
## [45]	train-rmse:0.324805+0.001692	train-auc:0.932658+0.001348	train-error:0.153733+0.002719	tes
## [46]	train-rmse:0.324063+0.001575	train-auc:0.933289+0.001216	train-error:0.153223+0.003100	tes
## [47]	train-rmse:0.323439+0.001762	train-auc:0.933793+0.001388	train-error:0.153135+0.003544	tes
## [48]	train-rmse:0.322854+0.001788	train-auc:0.934262+0.001377	train-error:0.152391+0.003528	tes
## [49]	train-rmse:0.322252+0.001745	train-auc:0.934807+0.001355	train-error:0.151750+0.003334	tes
## [50]	train-rmse:0.321600+0.001629	train-auc:0.935366+0.001298	train-error:0.151035+0.002842	tes
## [51]	train-rmse:0.320875+0.001747	train-auc:0.935970+0.001400	train-error:0.150423+0.003509	tes
## [52]	train-rmse:0.320162+0.001893	train-auc:0.936542+0.001521	train-error:0.149329+0.003485	tes
## [53]	train-rmse:0.319647+0.001986	train-auc:0.936932+0.001590	train-error:0.149081+0.003185	tes
## [54]	train-rmse:0.319180+0.002051	train-auc:0.937360+0.001665	train-error:0.148921+0.003133	tes
## [55]	train-rmse:0.318495+0.002023	train-auc:0.937930+0.001619	train-error:0.148046+0.002944	tes


```

## [56] train-rmse:0.317847+0.001816    train-auc:0.938420+0.001408    train-error:0.147171+0.002895    tes
## [57] train-rmse:0.317274+0.001728    train-auc:0.938898+0.001332    train-error:0.146530+0.002161    tes
## [58] train-rmse:0.316655+0.001743    train-auc:0.939373+0.001312    train-error:0.145669+0.002031    tes
## [59] train-rmse:0.316001+0.001710    train-auc:0.939908+0.001291    train-error:0.145217+0.002054    tes
## [60] train-rmse:0.315394+0.001733    train-auc:0.940395+0.001311    train-error:0.144736+0.002320    tes
## [61] train-rmse:0.314790+0.001697    train-auc:0.940824+0.001289    train-error:0.144430+0.002118    tes
## [62] train-rmse:0.314159+0.001587    train-auc:0.941314+0.001225    train-error:0.143992+0.001886    tes
## [63] train-rmse:0.313643+0.001755    train-auc:0.941697+0.001368    train-error:0.143336+0.002088    tes
## [64] train-rmse:0.313051+0.001694    train-auc:0.942164+0.001256    train-error:0.142695+0.002148    tes
## [65] train-rmse:0.312558+0.001828    train-auc:0.942522+0.001354    train-error:0.142695+0.002514    tes
## [66] train-rmse:0.312013+0.001876    train-auc:0.942916+0.001425    train-error:0.141951+0.002486    tes
## [67] train-rmse:0.311434+0.001698    train-auc:0.943358+0.001266    train-error:0.141441+0.002384    tes
## [68] train-rmse:0.310875+0.001640    train-auc:0.943777+0.001186    train-error:0.141134+0.002037    tes
## [69] train-rmse:0.310278+0.001617    train-auc:0.944223+0.001195    train-error:0.140245+0.001969    tes
## [70] train-rmse:0.309696+0.001594    train-auc:0.944673+0.001148    train-error:0.139531+0.002061    tes
## [71] train-rmse:0.309181+0.001532    train-auc:0.945069+0.001093    train-error:0.138947+0.001563    tes
## [72] train-rmse:0.308746+0.001576    train-auc:0.945353+0.001151    train-error:0.138320+0.001831    tes
## [73] train-rmse:0.308254+0.001624    train-auc:0.945715+0.001192    train-error:0.138087+0.001517    tes
## [74] train-rmse:0.307611+0.001518    train-auc:0.946229+0.001126    train-error:0.137270+0.001555    tes
## [75] train-rmse:0.307122+0.001547    train-auc:0.946569+0.001098    train-error:0.136468+0.001889    tes
## [76] train-rmse:0.306571+0.001594    train-auc:0.946961+0.001127    train-error:0.135944+0.001965    tes
## [77] train-rmse:0.305946+0.001564    train-auc:0.947416+0.001071    train-error:0.135185+0.001789    tes
## [78] train-rmse:0.305473+0.001577    train-auc:0.947754+0.001094    train-error:0.134835+0.002163    tes
## [79] train-rmse:0.305042+0.001672    train-auc:0.948028+0.001165    train-error:0.134354+0.002092    tes
## [80] train-rmse:0.304559+0.001754    train-auc:0.948360+0.001245    train-error:0.133596+0.002231    tes
## [81] train-rmse:0.304144+0.001801    train-auc:0.948627+0.001221    train-error:0.133231+0.002122    tes
## [82] train-rmse:0.303739+0.001734    train-auc:0.948914+0.001159    train-error:0.133348+0.002223    tes
## [83] train-rmse:0.303312+0.001882    train-auc:0.949232+0.001253    train-error:0.132458+0.002218    tes
## [84] train-rmse:0.302778+0.001926    train-auc:0.949588+0.001289    train-error:0.131715+0.002569    tes
## [85] train-rmse:0.302109+0.001988    train-auc:0.950044+0.001358    train-error:0.131030+0.002782    tes
## [86] train-rmse:0.301626+0.002012    train-auc:0.950331+0.001356    train-error:0.130767+0.002845    tes
## [87] train-rmse:0.300992+0.001877    train-auc:0.950830+0.001264    train-error:0.130300+0.002366    tes
## [88] train-rmse:0.300468+0.001871    train-auc:0.951155+0.001258    train-error:0.129863+0.002225    tes
## [89] train-rmse:0.299921+0.001991    train-auc:0.951500+0.001340    train-error:0.129251+0.002430    tes
## [90] train-rmse:0.299382+0.002110    train-auc:0.951856+0.001388    train-error:0.128653+0.002764    tes
## [91] train-rmse:0.298813+0.002120    train-auc:0.952226+0.001396    train-error:0.128055+0.002812    tes
## [92] train-rmse:0.298246+0.002226    train-auc:0.952622+0.001474    train-error:0.127647+0.002876    tes
## [93] train-rmse:0.297793+0.002160    train-auc:0.952924+0.001428    train-error:0.127238+0.002568    tes
## [94] train-rmse:0.297216+0.002172    train-auc:0.953288+0.001439    train-error:0.126947+0.002558    tes
## [95] train-rmse:0.296801+0.002170    train-auc:0.953534+0.001385    train-error:0.126465+0.002678    tes
## [96] train-rmse:0.296367+0.002323    train-auc:0.953836+0.001481    train-error:0.125722+0.002673    tes
## [97] train-rmse:0.295864+0.002369    train-auc:0.954187+0.001540    train-error:0.125882+0.002883    tes
## [98] train-rmse:0.295462+0.002359    train-auc:0.954423+0.001510    train-error:0.125299+0.002879    tes
## [99] train-rmse:0.295035+0.002460    train-auc:0.954663+0.001593    train-error:0.124993+0.002777    tes
## [100]    train-rmse:0.294489+0.002418    train-auc:0.955032+0.001500    train-error:0.124555+0.002784

```

```

trained_model <- xgb.train(data = xgb_train, max_depth = 3,
  eta = 1, nthread = 4, nrounds = 14,
  watchlist = list(train = xgb_train, eval = xgb_test),
  objective = "binary:logistic")

```

```

## [1] train-logloss:0.518785 eval-logloss:0.820399
## [2] train-logloss:0.460336 eval-logloss:0.881454
## [3] train-logloss:0.443383 eval-logloss:0.952815

```

```
## [4] train-logloss:0.429115 eval-logloss:1.006953
## [5] train-logloss:0.421462 eval-logloss:1.068660
## [6] train-logloss:0.417349 eval-logloss:1.082765
## [7] train-logloss:0.404430 eval-logloss:1.115529
## [8] train-logloss:0.401043 eval-logloss:1.130244
## [9] train-logloss:0.396576 eval-logloss:1.161850
## [10] train-logloss:0.393187 eval-logloss:1.185680
## [11] train-logloss:0.390200 eval-logloss:1.205786
## [12] train-logloss:0.388254 eval-logloss:1.226213
## [13] train-logloss:0.386829 eval-logloss:1.243061
## [14] train-logloss:0.383553 eval-logloss:1.241412
```

```
predictions <- predictions %>%
  mutate(Transported=ifelse(Transported>.5,'True','False'))

write_csv(predictions, "C:\\Users\\erico\\Desktop\\XGBpredictions_Rounds14.csv")
```

2. Tune SVM Radial: No Improvement

We manually hypertune our svm radial model by increasing sigma (loosening the fit) and decreasing sigma (tightening the fit). A decreased sigma increased accuracy on the holdout set but not on the Kaggle.

```
library(doParallel)
library(caret)

cl<-makePSOCKcluster(10)

registerDoParallel(cl)

#rad3 <- EHModel_SVM(dfTrain5zz, "Transported", method="radial", cValue=1, sigmaValue =.05)

rad3 <- EHModel_SVM(dfTrain5zz, "Transported", method="radial", cValue=1, sigmaValue =.04)
```

```
## Support Vector Machines with Radial Basis Function Kernel
##
## 6097 samples
## 27 predictor
## 2 classes: '0', '1'
##
## Pre-processing: centered (27), scaled (27)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 5487, 5487, 5487, 5488, 5488, 5488, ...
## Resampling results:
##
## Accuracy Kappa
## 0.7959085 0.5918032
##
## Tuning parameter 'sigma' was held constant at a value of 0.04
## Tuning
## parameter 'C' was held constant at a value of 1
## Confusion Matrix and Statistics
##
```

```

##           Reference
## Prediction    0    1
##           0 576 163
##           1 181 603
##
##           Accuracy : 0.7741
##           95% CI : (0.7523, 0.7949)
##           No Information Rate : 0.503
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.5482
##
## Mcnemar's Test P-Value : 0.3594
##
##           Sensitivity : 0.7609
##           Specificity : 0.7872
##           Pos Pred Value : 0.7794
##           Neg Pred Value : 0.7691
##           Prevalence : 0.4970
##           Detection Rate : 0.3782
##           Detection Prevalence : 0.4852
##           Balanced Accuracy : 0.7741
##
##           'Positive' Class : 0
##
predictions_rad3 <- EHModel_Predict(rad3$svm, dfTest5zz, predictionsColumnName = "Transported", testDat
predictions_rad3$Transported <- ifelse(predictions_rad3$Transported==1,"True","False")
write_csv(predictions_rad3, "C://Users//erico//Desktop//SVM_Rad_LHigherSigma.csv")
stopCluster(cl)

```

Discussion

(see above)